

EE2211 Introduction to Machine Learning

Lecture 1

Wang Xinchao
xinchao@nus.edu.sg

**Office Hour: Mon 9:30 – 10:30 AM
(Week 2-4, Week 10-12)**


Course Contents

- Introduction and Preliminaries (Xinchao)
 - **Introduction**
 - Data Engineering
 - Introduction to Probability and Statistics
- Fundamental Machine Learning Algorithms I (Yueming)
 - Systems of linear equations
 - Least squares, Linear regression
 - Ridge regression, Polynomial regression
- Fundamental Machine Learning Algorithms II (Yueming)
 - Over-fitting, bias/variance trade-off
 - Optimization, Gradient descent
 - Decision Trees, Random Forest
- Performance and More Algorithms (Xinchao)
 - Performance Issues
 - K-means Clustering
 - Neural Networks

World's Largest Selfie



World's Largest Selfie



Outline

- What is machine learning?
 - Three Definition(s)
- When do we need machine learning?
 - Sometimes we need, sometimes we don't
- Applications of machine learning
- Types of machine learning
 - Supervised, Unsupervised, Reinforcement Learning
- Walking through a toy example on classification
- Inductive vs. Deductive Reasoning

What is machine learning?

Learning is any process by which a system improves performance from experience.

- Herbert Simon


A computer program is said to learn

- from **experience E**
- with respect to some class of **tasks T**
- and **performance measure P**,

if its performance at tasks in T, as measured by P, improves with experience E.

- Tom Mitchell

Machine Learning (Supervised Learning)



Data Output



Cat

:

→ $f(\cdot)$ such that


Dog



$f(\text{Cat}) = \text{'cat'}$

$f(\text{Dog}) = \text{'dog'}$

Machine Learning (Supervised Learning)



Data Output



Cat

:



Dog

$$\longrightarrow f(\cdot)$$

When applied



$$f(\text{Cat}) \rightarrow \text{Cat !}$$

New image

Machine Learning: field of study that gives computers the ability to learn without being explicitly programmed

- Arthur Samuel

AI, Machine Learning, and Deep Learning

ARTIFICIAL INTELLIGENCE

Any technique which enables computers to mimic human behavior



1950's 1960's 1970's 1980's

MACHINE LEARNING

AI techniques that give computers the ability to learn without being explicitly programmed to do so



1990's 2000's 2010s

DEEP LEARNING

A subset of ML which make the computation of multi-layer neural networks feasible



Example of AI but not ML: Deductive Reasoning

NUS is in Singapore, Singapore is in Asia \rightarrow NUS is in Asia

When do we need machine learning?

Lack of human expertise
 (Navigating on Mars)



Involves huge amount of data
 (Genomics)



Learning is not always useful:

No need to “learn” to calculate payroll!

My Salary = Days_of_work * Daily Salary + Bonus

Application of Machine Learning

Task T, Performance P, Experience E

T: Digit Recognition


P: Classification Accuracy

E: Labelled Images

4 “four”

3 “three”

Labels -> Supervision!



Application of Machine Learning

Task T, Performance P, Experience E

T: Email Categorization

P: Classification Accuracy

E: Email Data, Some Labelled




Application of Machine Learning

Task T, Performance P, Experience E

T: Playing Go Game

P: Chances of Winning

E: Records of Past Games



Application of Machine Learning


Task T, Performance P, Experience E

T: Identifying Covid-19 Clusters

P: Small Internal Distances

Larger External Distances

E: Records of Patients





Web Search Engine



Product Recommendation



Language Translation



Photo Tagging




Virtual Personal Assistant



Portfolio Management



Traffic Prediction



Medical Diagnosis



Algorithmic Trading

Types of Machine Learning

Supervised Learning

Input:

- 1) Training Samples,
- 2) Desired Output
(Teacher/Supervision)

Output:

A rule that maps input to output

Unsupervised Learning

Input:

Samples

Output:

Underlying patterns in data

Reinforcement Learning

Input:

Sequence of States,
Actions, and
Delayed Rewards

Output:

Action Strategy: a rule
that maps the
environment to action

Types of Machine Learning

Supervised Learning

Input:

- 1) Training Samples,
- 2) Desired Output
(Teacher/Supervision)

Output:

A rule that maps input to output

Data Output



Cat

⋮

→ $f(\cdot)$ such that



Dog

Unsupervised Learning

Input:

Samples

Output:

Underlying patterns in data

Reinforcement Learning

Input:

Sequence of States,
Actions, and
Delayed Rewards

Output:


Action Strategy: a rule
that maps the
environment to action

$$f(\text{Cat}) = \text{'cat'}$$

$$f(\text{Dog}) = \text{'dog'}$$



Supervised Learning

Data




Regression

- Given $(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)$
- Learn a function $f(\mathbf{x})$ to predict real-valued y given \mathbf{x}



Supervised Learning



Data



Classification

- Given $(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)$
- Learn a function $f(\mathbf{x})$ to predict categorical y given \mathbf{x}

Feature Space



$f(x)$: line that separates two classes

Types of Machine Learning

Supervised Learning

Input:

- 1) Training Samples,
- 2) Desired Output
(Teacher/Supervision)

Output:

A rule that maps input to output

Unsupervised Learning

Input:

Samples

Output:

Underlying patterns in data

Reinforcement Learning

Input:

Sequence of States,
Actions, and
Delayed Rewards



Output:

Action Strategy: a rule
that maps the
environment to action

Unsupervised Learning

Clustering

- Given $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$, without labels
- Output Hidden Structure Behind



No Label/Supervision is given!

Types of Machine Learning

Supervised Learning

Input:

- 1) Training Samples,
- 2) Desired Output
(Teacher/Supervision)

Output:

A rule that maps input to output

Unsupervised Learning

Input:

Samples

Output:

Underlying patterns in data

Reinforcement Learning

Input:


Sequence of States,
Actions, and
Delayed Rewards

Output:


Action Strategy: a rule
that maps the
environment to action

Reinforcement Learning


Breakout Game



Initial Performance





Training 15 minutes



Training 30 minutes

Reinforcement Learning

- Given sequence of states S and actions A with (delayed) rewards R
- Output a policy $\pi(a, s)$, to guide us what action a to take in state s



S : Ball Location,
Paddle Location, Bricks

A : left, right

R :
positive reward
 Knocking a brick,
 clearing all bricks

negative reward
 Missing the ball

zero reward
 Cases in between

Supervised
Unsupervised
Reinforcement

Quiz Time!

0 0 0 1 1 1 1 1 2

2 2 2 2 2 2 3 3 3

3 4 4 4 4 5 5 5


6 6 6 7 7 7 7 8 8 8

9 9 9 9 9 9 9 9


Supervised



Supervised




Unsupervised




Reinforcement


Walking Through A Toy Example: Token Classification



Yes



No



Step 1: Feature Extraction

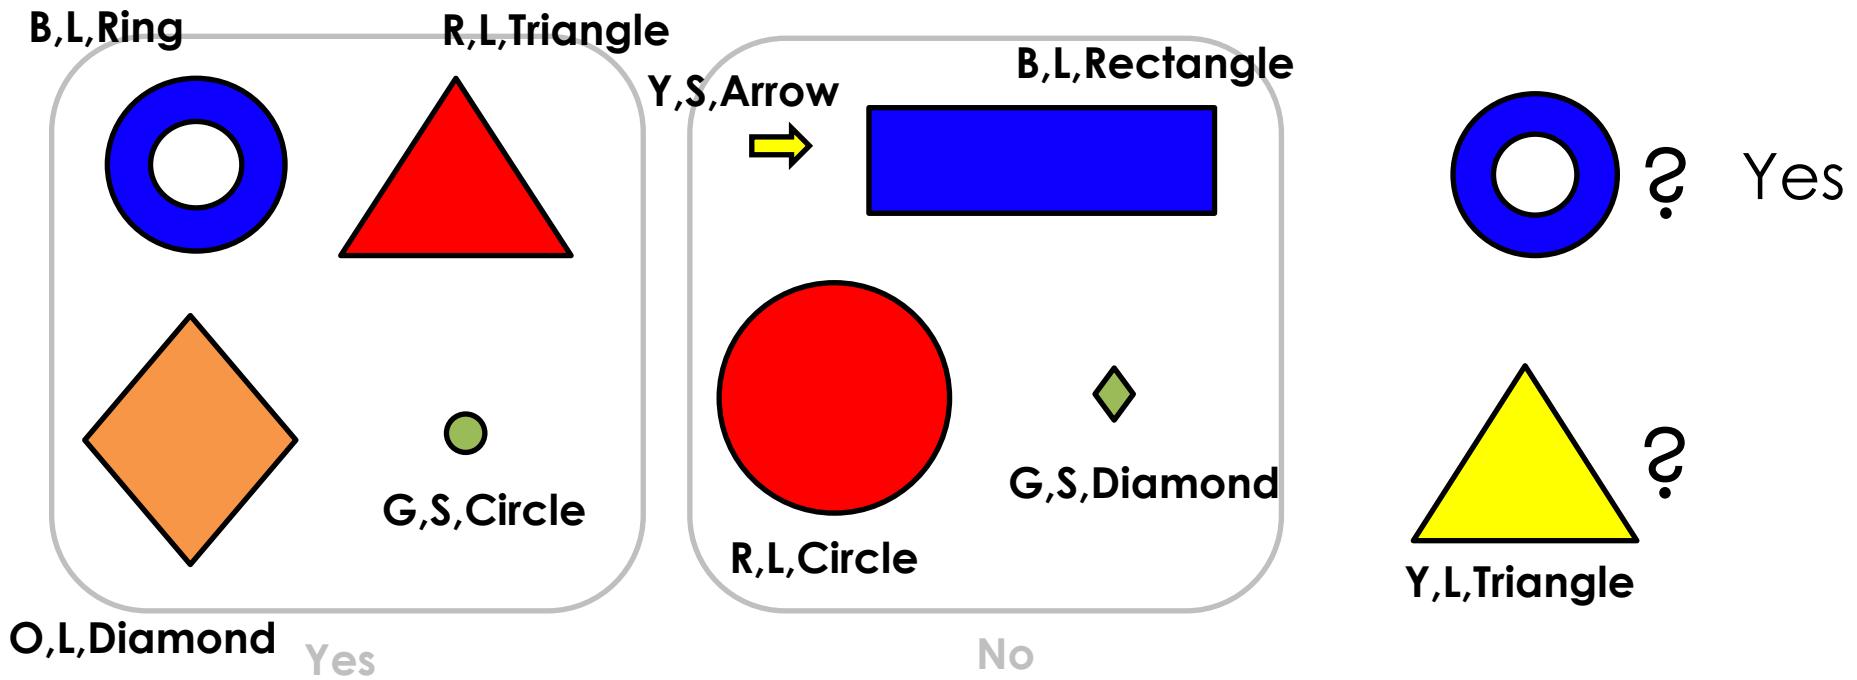
Extract Attributes of Samples



Step 2: Sample Classification

Decide Label for a Sample

Walking Through A Toy Example: Token Classification



Step 1: Feature Extraction

Color, Size, Shape

Walking Through A Toy Example: Token Classification

Feature Extraction

	Color	Size	Shape	Label
O	Blue	Large	Ring	Yes
▲	Red	Large	Triangle	Yes
◆	Orange	Large	Diamond	Yes
●	Green	Small	Circle	Yes
→	Yellow	Small	Arrow	No
■	Blue	Large	Rectangle	No
●	Red	Large	Circle	No
◆	Green	Small	Diamond	No
▲	Yellow	Large	Triangle	?

Walking Through A Toy Example: Token Classification

Feature Extraction

	Color	Size	Shape	Label
O	Blue	Large	Ring	Yes
▲	Red	Large	Triangle	Yes
◆	Orange	Large	Diamond	Yes
●	Green	Small	Circle	Yes
→	Yellow	Small	Arrow	No
■	Blue	Large	Rectangle	No
○	Red	Large	Circle	No
◆	Green	Small	Diamond	No

Walking Through A Toy Example: Token Classification

Feature Extraction

Color	Size	Shape	Label
Blue	Large	Ring	Yes
Red	Large	Triangle	Yes
Orange	Large	Diamond	Yes
Green	Small	Circle	Yes
Yellow	Small	Arrow	No
Blue	Large	Rectangle	No
Red	Large	Circle	No
Green	Small	Diamond	No


Similarity

Color	Size	Shape	Total
0	1	0	1
0	1	1	2
0	1	0	1
0	0	0	0
1	0	0	1
0	1	0	1
0	1	0	1
0	0	0	0

Legend:

- Yellow triangle: 0
- Red triangle: 1
- Orange diamond: 0
- Green circle: 0
- Yellow arrow: 1
- Blue rectangle: 1
- Red circle: 1
- Green diamond: 0

Walking Through A Toy Example: Token Classification



Similarity


	Color	Size	Shape	Total
▲ ○	0	1	0	1
▲ ▲	0	1	1	2
▲ ◊	0	1	0	1
▲ ●	0	0	0	0
▲ →	1	0	0	1
▲ ■	0	1	0	1
▲ ○	0	1	0	1
▲ ♦	0	0	0	0

Nearest Neighbor Classifier:


- 1) Find the “nearest neighbor” of a sample in the feature space
- 2) Assign the label of the nearest neighbor to the sample

Inductive vs. Deductive Reasoning

- Main Task of Machine Learning: to make inference



Probability and Statistics



Rule-based reasoning

NUS is in Singapore, Singapore is in Asia =>
NUS is in Asia

Inductive Reasoning

Note: humans use inductive reasoning all the time and not in a formal way like using probability/statistics.

B. C.

by Johnny hart



Ref: Gardner, Martin (March 1979). "[MATHEMATICAL GAMES: On the fabric of inductive logic, and some probability paradoxes](#)" (PDF). *Scientific American*. 234

Summary by Quick Quiz

Three Components in ML Definition

Task T, Performance P, Experience E

Three Types of in ML

Supervised Learning
Unsupervised Learning
Reinforcement Learning

Two Types of Supervised Learning

Classification, Regression

One Type of Unsupervised Learning

Clustering

Inductive and Deductive

Inductive: Probable
Deductive: Rule-based

Example of a Classifier Model

Nearest Neighbor Classifier

