

EE2211 Introduction to Machine Learning

Lecture 1

Wang Xinchao
xinchao@nus.edu.sg

**Office Hour: Mon 9:30 – 10:30 AM
(Week 2-4, Week 10-12)**

Course Contents

1

- Introduction and Preliminaries (Xinchao)

- Introduction
- Data Engineering
- Introduction to Probability and Statistics

*No coding for
Exams*

- Fundamental Machine Learning Algorithms I (Yueming)

2

- Systems of linear equations
- Least squares, Linear regression
- Ridge regression, Polynomial regression

- Fundamental Machine Learning Algorithms II (Yueming)

3

- Over-fitting, bias/variance trade-off
- Optimization, Gradient descent
- Decision Trees, Random Forest

- Performance and More Algorithms (Xinchao)

4

- Performance Issues
- K-means Clustering
- Neural Networks

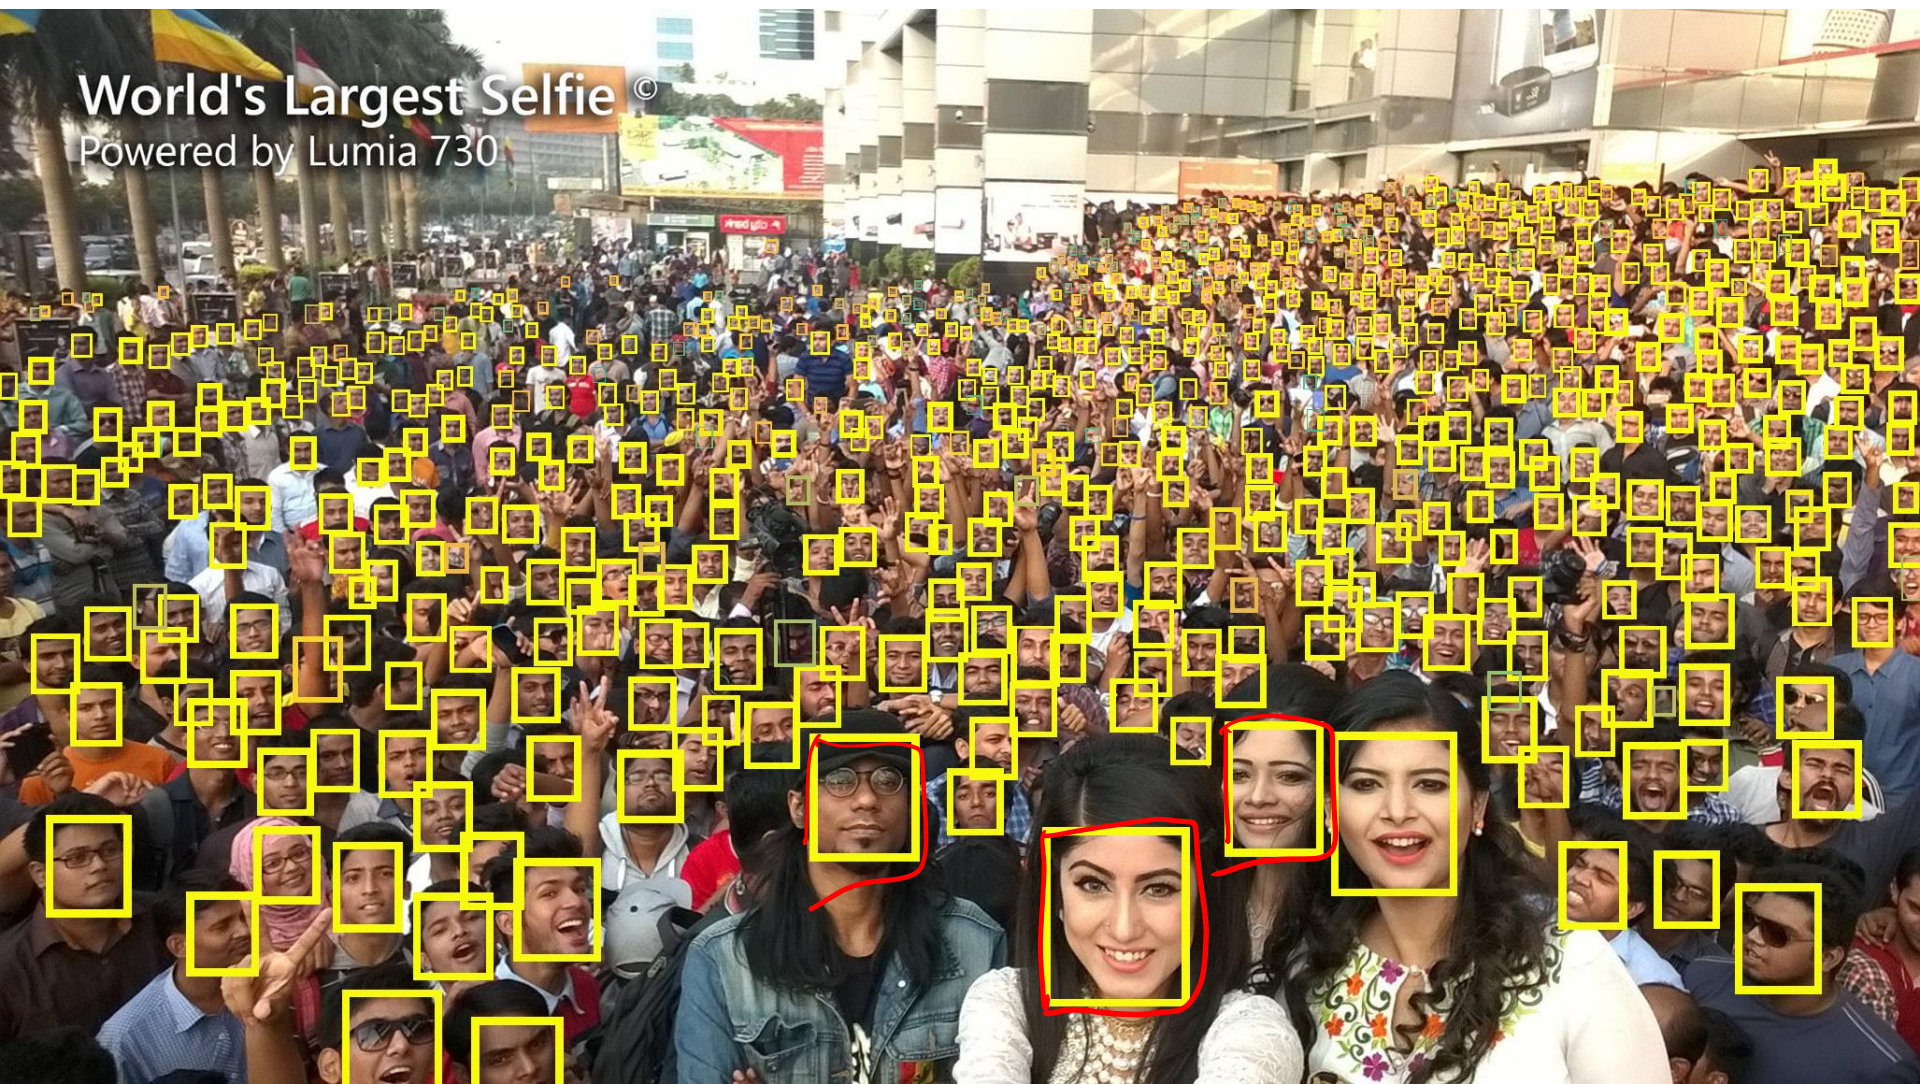
*No coding for
Exams*

World's Largest Selfie

World's Largest Selfie ©
Powered by Lumia 730



World's Largest Selfie



Outline

- What is machine learning?
 - Three Definition(s)
- When do we need machine learning?
 - Sometimes we need, sometimes we don't
- Applications of machine learning
- Types of machine learning
 - Supervised, Unsupervised, Reinforcement Learning
- Walking through a toy example on classification
- Inductive vs. Deductive Reasoning

What is machine learning?

Learning is any process by which a system improves performance from experience.↑

- Herbert Simon

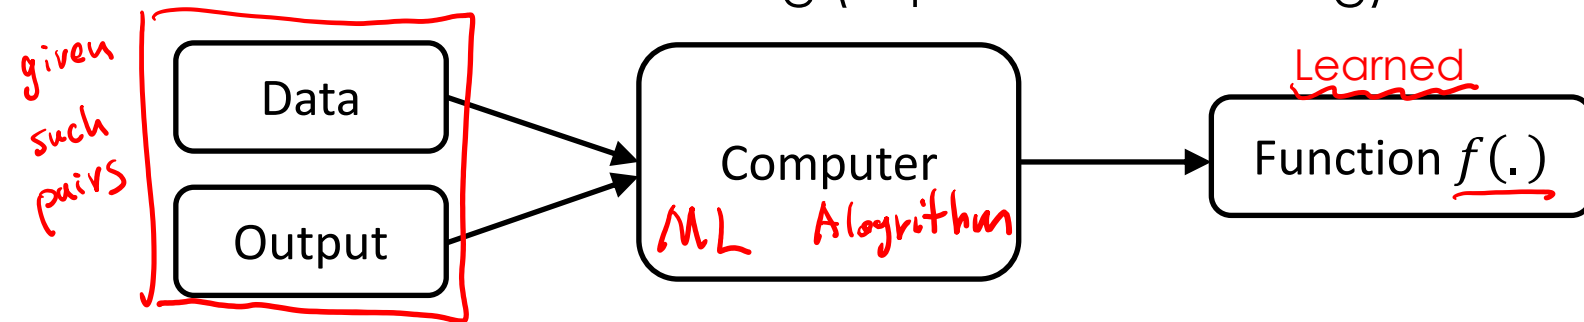
A computer program is said to learn

- from **experience E**
- with respect to some class of **tasks T**
- and **performance measure P** ,


if its performance at tasks in T , as measured by P , improves with experience E .

- Tom Mitchell

Machine Learning (Supervised Learning)



Data Output / label



→ Cat

⋮


→ $f(.)$ such that




→ Dog

$$f(\text{cat image}) = \text{'cat'}$$
$$f(\text{dog image}) = \text{'dog'}$$

Machine Learning (Supervised Learning)



Data **Output**



Cat

⋮

→ $f(.)$



Dog

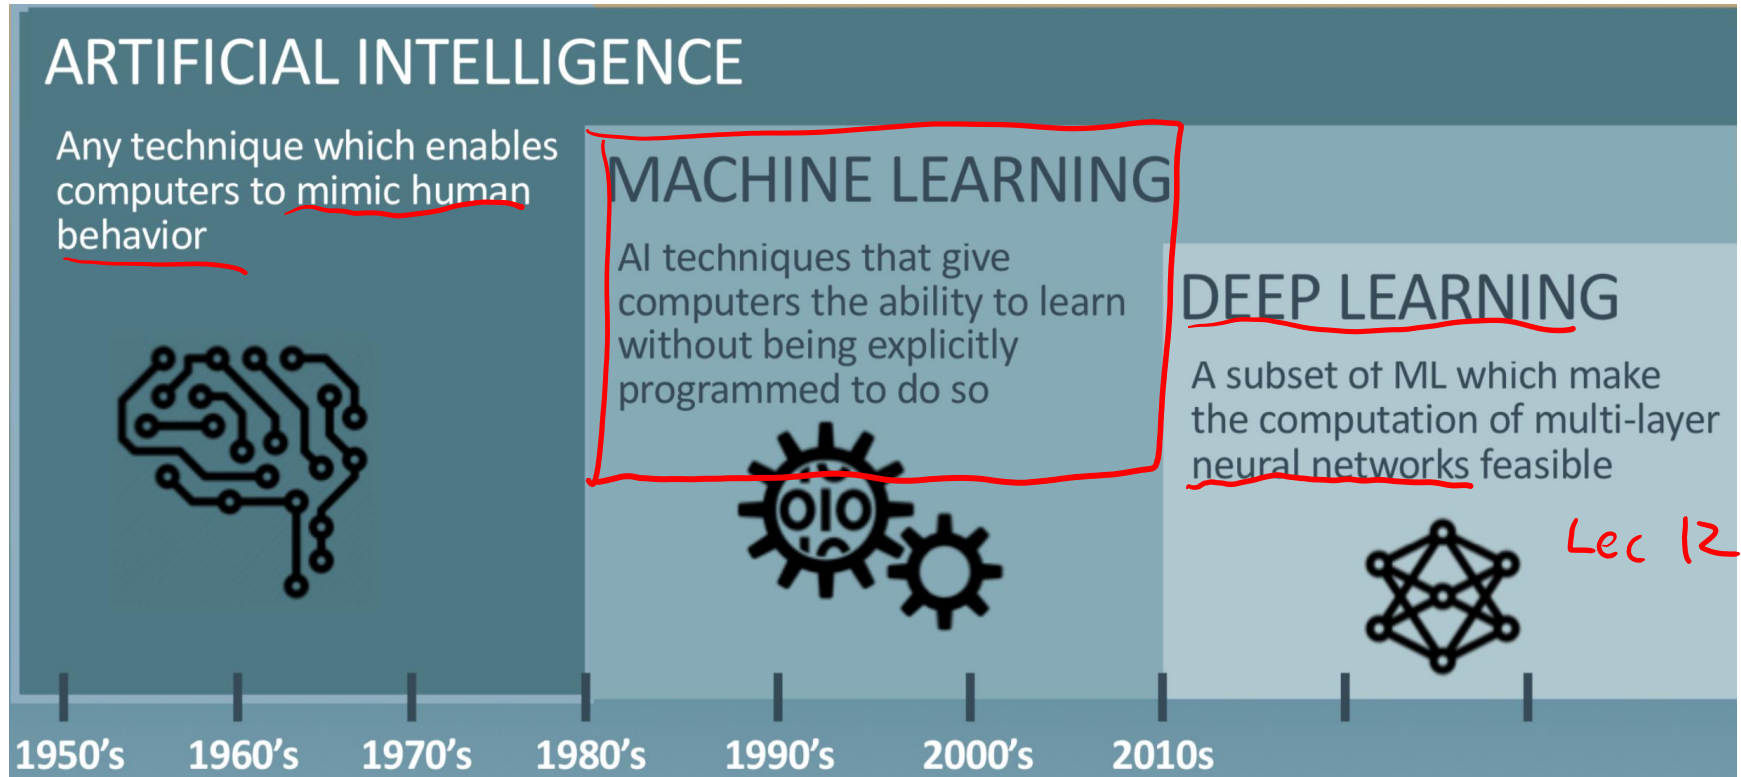
When applied

$f(\text{New image}) \rightarrow \text{Cat !}$

Machine Learning: field of study that gives computers the ability to learn without being explicitly programmed

- Arthur Samuel

AI, Machine Learning, and Deep Learning



Example of AI but not ML: Deductive Reasoning

NUS is in Singapore, Singapore is in Asia -> NUS is in Asia

When do we need machine learning?

Lack of human expertise
(Navigating on Mars)



Involves huge amount of data
(Genomics)



Learning is not always useful:

No need to “learn” to calculate payroll!

$$\text{My Salary} = \text{Days_of_work} * \text{Daily Salary} + \text{Bonus}$$

3,500
 30
 100
 500

Application of Machine Learning

Task T, Performance P, Experience E

VSPS

T: Digit Recognition

P: Classification Accuracy

E: Labelled Images

0 0 0 1 1 1 1 1 1 2

2 2 2 2 2 2 2 3 3 3

3 4 4 4 4 4 5 5 5 5

6 6 7 7 7 7 8 8 8 8

9 9 9 9 9 9 9 9 9 9

4 "four"

3 "three"

Labels -> Supervision!


Application of Machine Learning

Task T, Performance P, Experience E

T: Email Categorization

P: Classification Accuracy

E: Email Data, Some Labelled



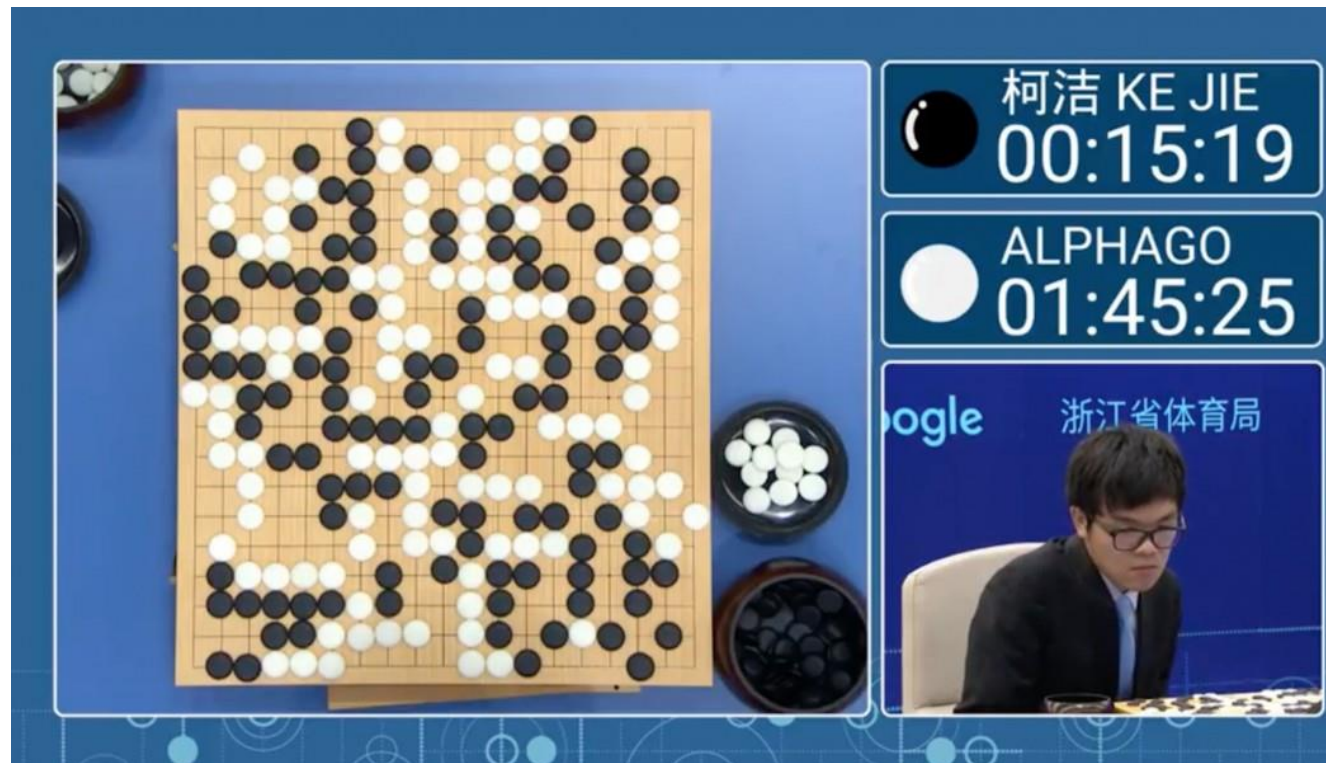
Application of Machine Learning

Task T, Performance P, Experience E

T: Playing Go Game

P: Chances of Winning

E: Records of Past Games



Application of Machine Learning

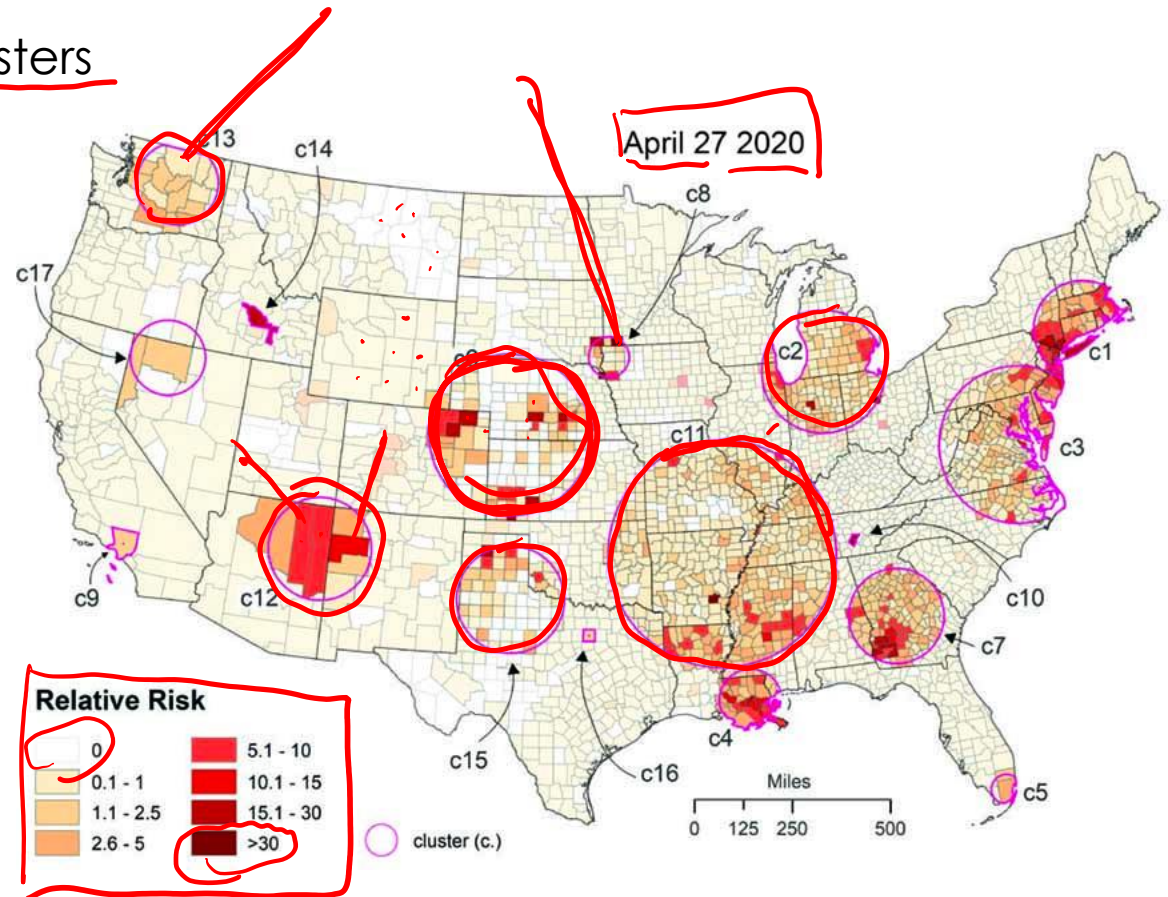
Task T, Performance P, Experience E

T: Identifying Covid-19 Clusters

P: Small Internal Distances

Larger External Distances

E: Records of Patients





Web Search Engine



Product Recommendation



Language Translation




Photo Tagging




Virtual Personal Assistant



Portfolio Management



Traffic Prediction



Medical Diagnosis



Algorithmic Trading

Types of Machine Learning

Supervised Learning

Input:

- 1) Training Samples,
- 2) Desired Output (Teacher/Supervision)

Output:

A rule that maps input to output

Unsupervised Learning

Input:

Samples

Output:

Underlying patterns in data

Reinforcement Learning

Input:

Sequence of States, Actions, and Delayed Rewards

Output:

Action Strategy: a rule that maps the environment to action

Types of Machine Learning

Supervised Learning


Input:

- 1) Training Samples,
- 2) Desired Output (Teacher/Supervision)

Output: *function*

A rule that maps input to output

Data Output




Cat

⋮



Dog



Unsupervised Learning

Input:

Samples

Output:

Underlying patterns in data

$$f(\text{cat image}) = \text{'cat'}$$

$$f(\text{dog image}) = \text{'dog'}$$

Reinforcement Learning

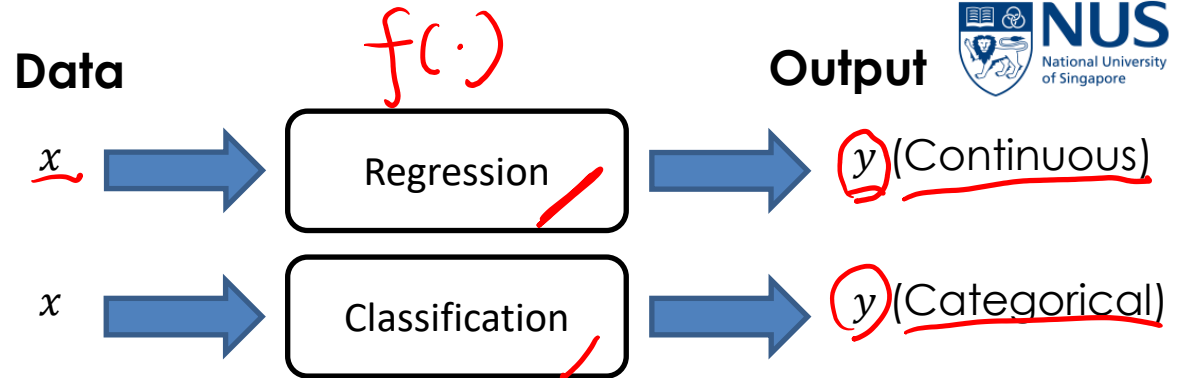
Input:

Sequence of States, Actions, and Delayed Rewards

Output:

Action Strategy: a rule that maps the environment to action



Supervised Learning



Regression

- Given $(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)$
- Learn a function $f(x)$ to predict real-valued y given x

linear regression



Supervised Learning

Data

x



Regression



Output

y (Continuous)

x




Classification



y (Categorical)

Classification

- Given $(\underline{x}_1, y_1), (\underline{x}_2, y_2), \dots, (\underline{x}_N, y_N)$
- Learn a function $f(\underline{x})$ to predict categorical y given \underline{x}



Types of Machine Learning

Supervised Learning

Input:

- 1) Training Samples,
- 2) ~~Desired Output~~
(Teacher/Supervision)

Output:

A rule that maps input to output

Unsupervised Learning

Input:

Samples

Output:

Underlying patterns in data

Reinforcement Learning

Input:

Sequence of States,
Actions, and
Delayed Rewards

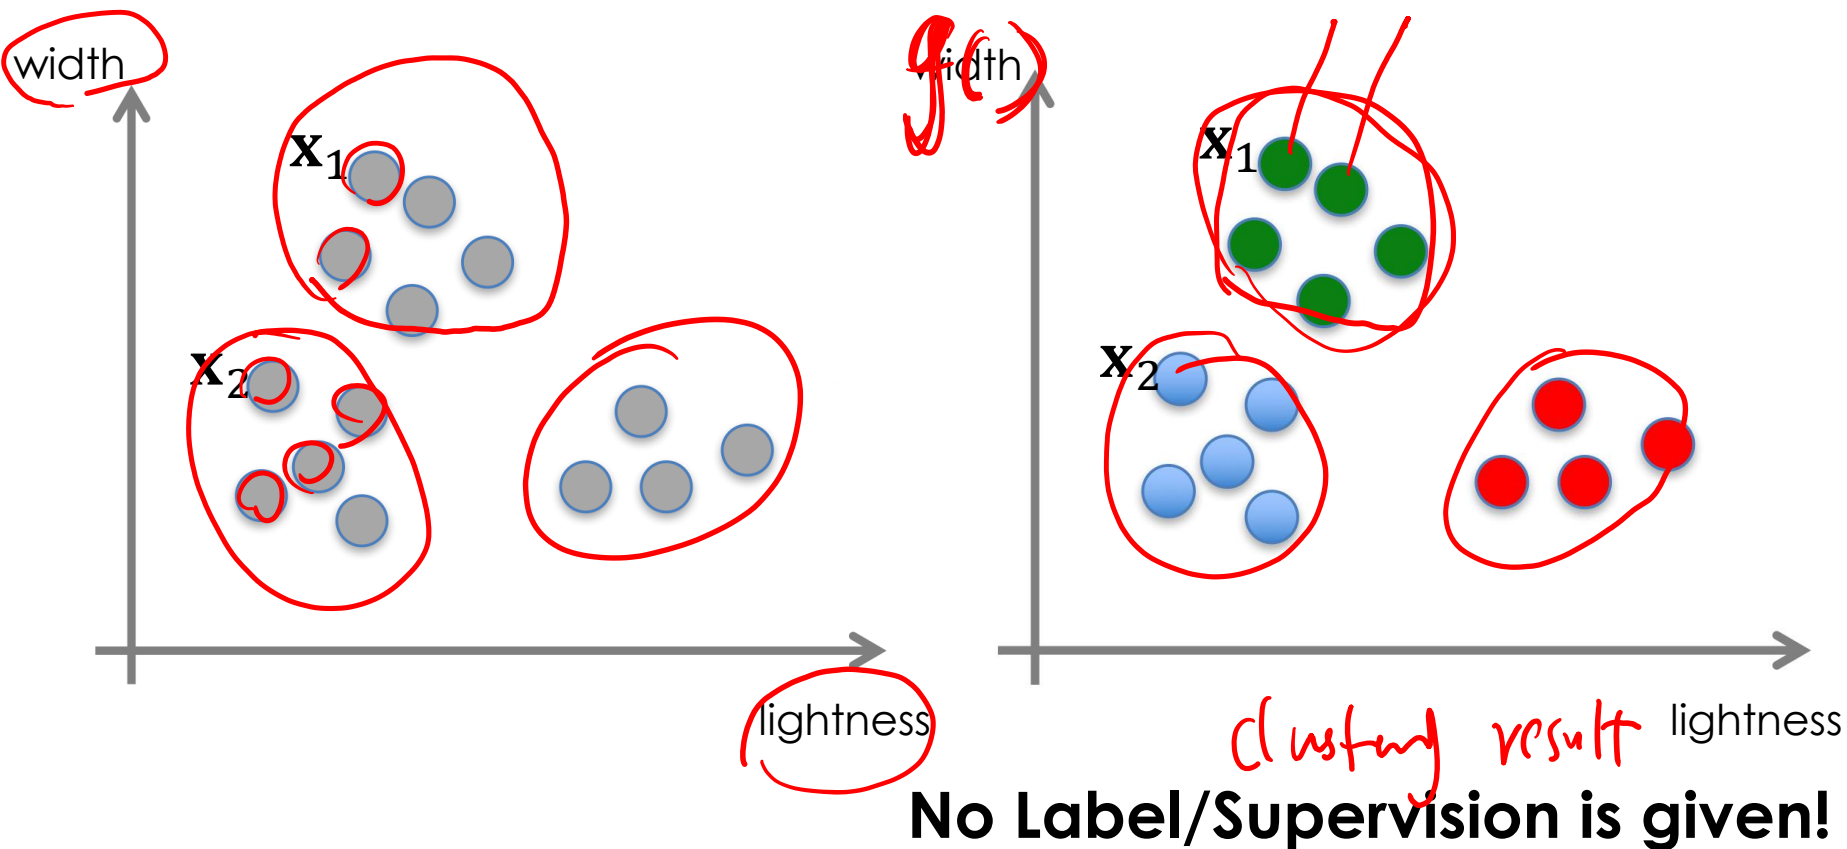
Output:

Action Strategy: a rule
that maps the
environment to action

Unsupervised Learning

Clustering

- Given $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$, without labels
- Output Hidden Structure Behind



Types of Machine Learning

Supervised Learning

Input:

- 1) Training Samples,
- 2) Desired Output (Teacher/Supervision)

Output:

A rule that maps input to output

Unsupervised Learning

Input:

Samples

Output:

Underlying patterns in data

advanced topic

Reinforcement Learning

Input:


Sequence of States,
Actions, and
Delayed Rewards

Output:


Action Strategy: a rule
that maps the
environment to action

Reinforcement Learning


Breakout Game



Initial Performance





Training 15 minutes



Training 30 minutes

Reinforcement Learning

- Given sequence of states S and actions A with (delayed) rewards R
- Output a policy $\pi(a, s)$, to guide us what action a to take in state s



S: Ball Location,
Paddle Location, Bricks

A: left, right


R:
positive reward
Knocking a brick,
clearing all bricks

negative reward
Missing the ball

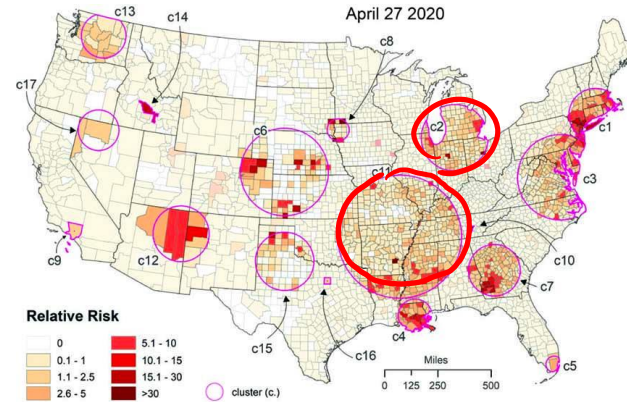
zero reward
Cases in between

Supervised
Unsupervised
Reinforcement

Quiz Time!



Supervised



Unsupervised




Supervised



Reinforcement

Walking Through A Toy Example: Token Classification



Step1: Feature Extraction

Extract Attributes of Samples




Step2: Sample Classification

Decide Label for a Sample

Nearest Neighbor Classifier

Walking Through A Toy Example: Token Classification












Step 1: Feature Extraction

Color, Size, Shape

Feature Selection

Walking Through A Toy Example: Token Classification

Feature Extraction

	Color —	Size —	Shape —	Label ·
	Blue	Large	Ring	Yes
	Red	Large	Triangle	Yes
	Orange	Large	Diamond	Yes
	Green	Small	Circle	Yes
	Yellow	Small	Arrow	No
	Blue	Large	Rectangle	No
	Red	Large	Circle	No
	Green	Small	Diamond	No
	Yellow	Large	Triangle	?

Walking Through A Toy Example: Token Classification

Feature Extraction









	Color	Size	Shape	Label
	Blue	Large	Ring	Yes
	Red	Large	Triangle	Yes
	Orange	Large	Diamond	Yes
	Green	Small	Circle	Yes
	Yellow	Small	Arrow	No
	Blue	Large	Rectangle	No
	Red	Large	Circle	No
	Green	Small	Diamond	No

Walking Through A Toy Example: Token Classification

Feature Extraction

	Color	Size	Shape	Label
	Blue	Large	Ring	Yes
	Red	Large	Triangle	Yes
	Orange	Large	Diamond	Yes
	Green	Small	Circle	Yes
	Yellow	Small	Arrow	No
	Blue	Large	Rectangle	No
	Red	Large	Circle	No
	Green	Small	Diamond	No

















Similarity

	Color	Size	Shape	Total
	0	1	0	1
	0	1	1	2
	0	1	0	1
	0	0	0	0
	1	0	0	1
	0	1	0	1
	0	1	0	1
	0	0	0	0

Walking Through A Toy Example: Token Classification

'Yes'

Similarity

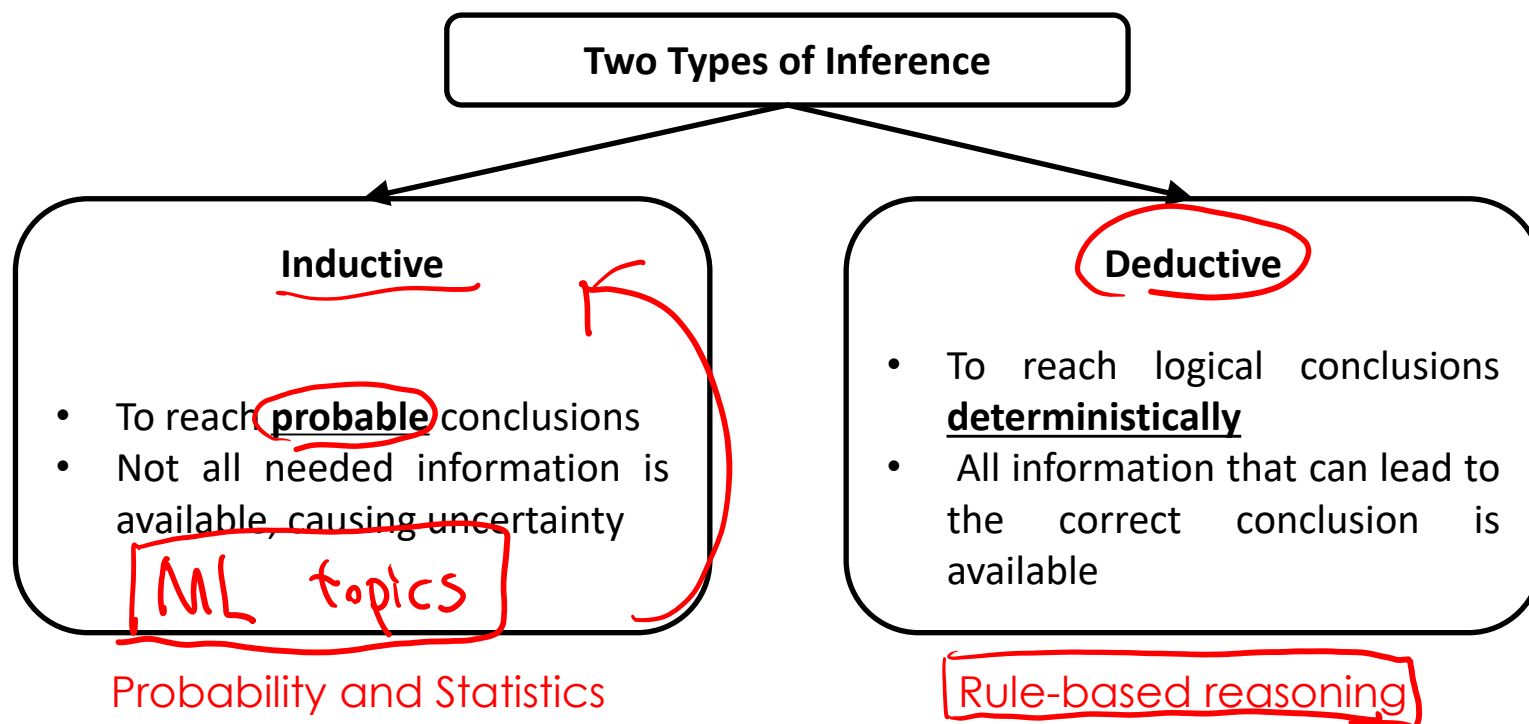
	Color ✓	Size ✓	Shape ✓	Total
 	0	1	0	1
 	0	1	1	2
 	0	1	0	1
 	0	0	0	0
 	1	0	0	1
 	0	1	0	1
 	0	1	0	1
 	0	0	0	0

Nearest Neighbor Classifier:


- 1) Find the "nearest neighbor" of a sample in the feature space
- 2) Assign the label of the nearest neighbor to the sample

Inductive vs. Deductive Reasoning

- Main Task of Machine Learning: to make inference



Probability and Statistics



NUS is in Singapore, Singapore is in Asia ->
NUS is in Asia

Inductive Reasoning

Note: humans use ~~inductive~~ reasoning all the time and not in a formal way like using probability/statistics.

B. C.

by Johnny hart



Ref: Gardener, Martin (March 1979). "MATHEMATICAL GAMES: On the fabric of inductive logic, and some probability paradoxes" (PDF). *Scientific American*. 234

Summary by Quick Quiz

Three Components in ML Definition

Task T, Performance P, Experience E

Three Types of in ML

Supervised Learning

Unsupervised Learning

Reinforcement Learning

Two Types of Supervised Learning

Classification, Regression

One Type of Unsupervised Learning

Clustering

Inductive and Deductive

Inductive: Probable

Deductive: Rule-based

Example of a Classifier Model

Nearest Neighbor Classifier

Practice Question

(Type of Question to Expect in Exams)

Which of the following statement is true?

- A. Nearest Neighbor Classifier is an example of supervised learning
- B. Nearest Neighbor Classifier is an example of deductive learning
- C. Nearest Neighbor Classifier is an example of feature extraction
- D. None of the above is correct.

