

EE2211 Introduction to Machine Learning

Lecture 5
Semester 2
2023/2024

Yueming Jin
ymjin@nus.edu.sg

Electrical and Computer Engineering Department
National University of Singapore

Acknowledgement:
EE2211 development team
(Xinchao, Helen, Thomas, Kar-Ann, Vincent, Chen Khong, Robby, and Haizhou)

Course Contents

- Introduction and Preliminaries (Xinchao)
 - Introduction
 - Data Engineering
 - Introduction to Probability and Statistics
- Fundamental Machine Learning Algorithms I (Yueming)
 - Systems of linear equations
 - **Least squares, Linear regression**
 - Ridge regression, Polynomial regression
- Fundamental Machine Learning Algorithms II (Yueming)
 - Over-fitting, bias/variance trade-off
 - Optimization, Gradient descent
 - Decision Trees, Random Forest
- Performance and More Algorithms (Xinchao)
 - Performance Issues
 - K-means Clustering
 - Neural Networks

Least Squares and Linear Regression

Module II Contents

- Notations, Vectors, Matrices (introduced in L3)
- Operations on Vectors and Matrices
- Systems of Linear Equations
- Set and Functions
- **Derivative and Gradient**
- **Least Squares, Linear Regression**
- **Linear Regression with Multiple Outputs**
- Linear Regression for Classification
- Ridge Regression
- Polynomial Regression

Recap: Linear and Affine Functions

Linear Functions

A function $f: \mathcal{R}^d \rightarrow \mathcal{R}$ is **linear** if it satisfies the following **two properties**:

- **Homogeneity** $f(\underline{\alpha \mathbf{x}}) = \underline{\alpha f(\mathbf{x})}$ **Scaling**
- **Additivity** $f(\underline{\mathbf{x} + \mathbf{y}}) = f(\underline{\mathbf{x}}) + f(\underline{\mathbf{y}})$ **Adding**

Inner product function

$$f(\mathbf{x}) = \mathbf{a}^T \mathbf{x} = a_1 x_1 + a_2 x_2 + \cdots a_d x_d$$

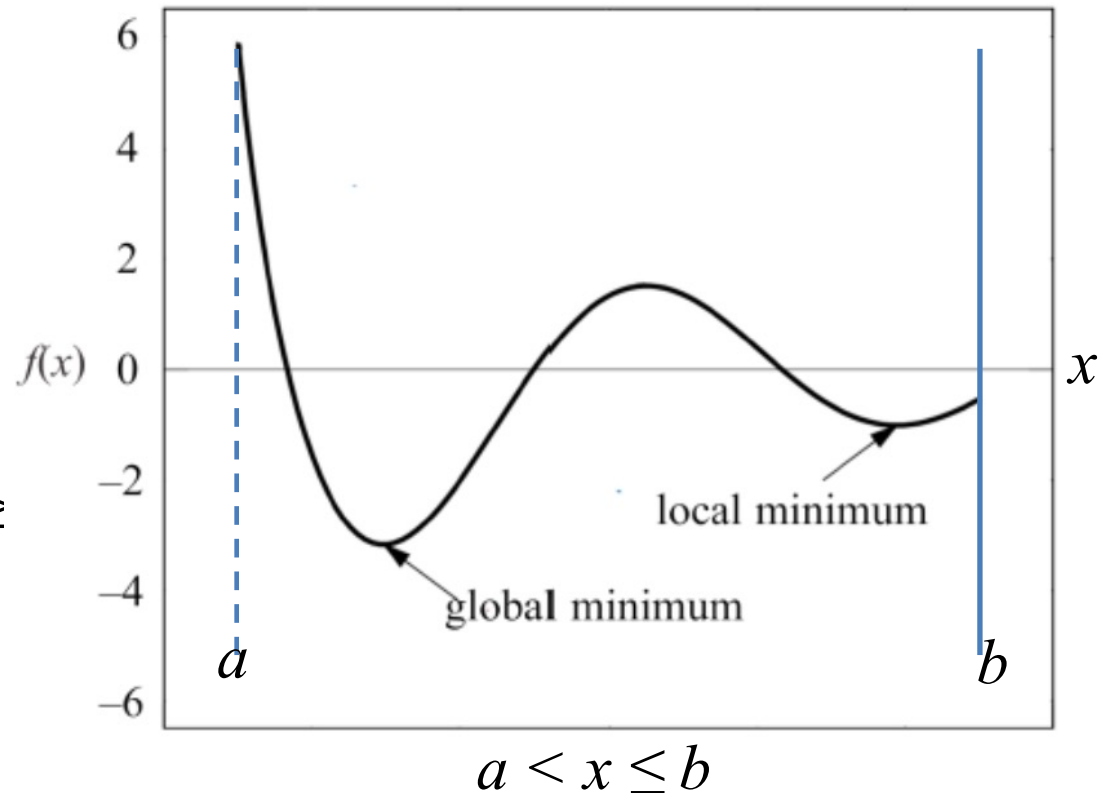
Affine function

$$f(\mathbf{x}) = \mathbf{a}^T \mathbf{x} + b \quad \text{scalar } b \text{ is called the offset (or bias)}$$

Functions: Maximum and Minimum

A local and a global minima of a function

- $f(x)$ has a **local minimum** at $x = c$ if $f(x) \geq f(c)$ for every x in some open interval around $x = c$
- $f(x)$ has a **global minimum** at $x = c$ if $f(x) \geq f(c)$ for all x in the domain of f



Note: An **interval** is a set of real numbers with the property that any number that lies between two numbers in the set is also included in the set.

An **open interval** does not include its endpoints and is denoted using parentheses. E.g. $(0, 1)$ means “all numbers greater than 0 and less than 1”.

Ref: [Book1] Andriy Burkov, “The Hundred-Page Machine Learning Book”, 2019 (p6-7 of chp2).

Functions: Maximum and Minimum

Max and Arg Max

- Given a set of values $\mathcal{A} = \{a_1, a_2, \dots, a_m\}$,
- The operator $\max_{a \in \mathcal{A}} f(a)$ returns the highest value $f(a)$ for all elements in the set \mathcal{A}
- The operator $\arg \max_{a \in \mathcal{A}} f(a)$ returns the element of the set \mathcal{A} that maximizes $f(a)$
- When the set is **implicit** or **infinite**, we can write

$$\max_a f(a) \quad \text{or} \quad \arg \max_a f(a)$$

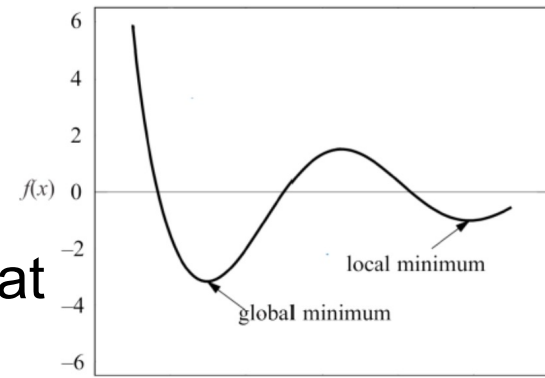
E.g. $f(a) = 3a, a \in [0,1] \rightarrow \max_a f(a) = 3$ and $\arg \max_a f(a) = 1$

Min and **Arg Min** operate in a similar manner

Note: **arg max** returns a value from the **domain** of the function and **max** returns from the **range (codomain)** of the function.

Derivative and Gradient

- The **derivative** f' of a function f is a function that describes how fast f grows (or decreases)
 - If the derivative is a **constant** value, e.g. 5 or -3
 - The function f grows (or decreases) constantly at any point x of its domain
 - When the derivative f' is a function
 - If f' is **positive at some x** , then the function f **grows** at this point
 - If f' is **negative at some x** , then the function f **decreases** at this point
 - The derivative of **zero at x** means that the function's **slope** at x is **horizontal** (e.g. maximum or minimum points)
- The process of finding a derivative is called **differentiation**.
- **Gradient** is the generalization of derivative for functions that take several inputs (or one input in the form of a vector or some other complex structure).



Ref: [Book1] Andriy Burkov, "The Hundred-Page Machine Learning Book", 2019 (p8 of chp2).

Derivative and Gradient

The gradient of a function is a vector of **partial derivatives**

Differentiation of a **scalar** function w.r.t. a **vector**

If $f(\mathbf{x})$ is a **scalar function** of d variables, \mathbf{x} is a $d \times 1$ vector.
Then differentiation of $f(\mathbf{x})$ w.r.t. \mathbf{x} results in a $d \times 1$ vector

$$\frac{df(\mathbf{x})}{d\mathbf{x}} = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \vdots \\ \frac{\partial f}{\partial x_d} \end{bmatrix}$$

$$\begin{array}{c} \text{I} : \mathbb{R}^3 \rightarrow \mathbb{R}^1 \\ (\underline{x}, \underline{y}, \underline{z}) \\ 3 \times 1 \end{array}$$

This is referred to as the **gradient** of $f(\mathbf{x})$ and often written as $\nabla_{\mathbf{x}} f$.

$$\text{E.g. } f(\mathbf{x}) = ax_1 + bx_2 \quad \nabla_{\mathbf{x}} f = \begin{bmatrix} a \\ b \end{bmatrix}$$

Ref: Duda, Hart, and Stork, "Pattern Classification", 2001 (Appendix)

Derivative and Gradient

Partial Derivatives

Differentiation of a **vector** function w.r.t. a **vector**

If $\mathbf{f}(\mathbf{x})$ is a **vector function** of size $h \times 1$ and \mathbf{x} is a $d \times 1$ vector.
Then differentiation of $\mathbf{f}(\mathbf{x})$ results in a $h \times d$ matrix

$$\frac{d\mathbf{f}(\mathbf{x})}{d\mathbf{x}} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_d} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_h}{\partial x_1} & \cdots & \frac{\partial f_h}{\partial x_d} \end{bmatrix}$$

The matrix is referred to as the **Jacobian** of $\mathbf{f}(\mathbf{x})$

Derivative and Gradient

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}$$

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \end{bmatrix}$$

Some Vector-Matrix Differentiation Formulae

$$\frac{d(\mathbf{Ax})}{d\mathbf{x}} = \mathbf{A}$$

Handwritten annotations: 2×1 above $d(\mathbf{Ax})$, 3×1 below $d\mathbf{x}$, and 2×3 below \mathbf{A} . A red arrow points from $d(\mathbf{Ax})$ to the expression:

$$\begin{bmatrix} a_{11}x_1 + a_{12}x_2 + a_{13}x_3 \\ a_{21}x_1 + a_{22}x_2 + a_{23}x_3 \end{bmatrix}$$

$$\frac{d(\mathbf{b}^T \mathbf{x})}{d\mathbf{x}} = \mathbf{b}$$

$$\frac{d(\mathbf{y}^T \mathbf{Ax})}{d\mathbf{x}} = \mathbf{A}^T \mathbf{y}$$

Handwritten annotations: \mathbf{b} is circled in red in the first equation. In the second equation, $\mathbf{y}^T \mathbf{A}$ is boxed in red, and \mathbf{b}^T is written above the box. $\mathbf{A}^T \mathbf{y}$ is underlined in red.

$$\frac{d(\mathbf{x}^T \mathbf{Ax})}{d\mathbf{x}} = (\mathbf{A} + \mathbf{A}^T) \mathbf{x}$$

$$f(\mathbf{x}) = \mathbf{a}^T \mathbf{x} = \underline{a_1 x_1 + a_2 x_2 + \dots + a_d x_d}$$

Derivations: <https://www.math.uwaterloo.ca/~hwolkowi/matrixcookbook.pdf>

Ref: Duda, Hart, and Stork, "Pattern Classification", 2001 (Appendix)

Linear Regression

- **Linear regression** is a popular regression learning algorithm that learns a model which is a linear combination of features of the input example.

$$\mathbf{X}\mathbf{w} = \mathbf{y}, \quad \mathbf{X} \in \mathcal{R}^{m \times d}, \quad \mathbf{w} \in \mathcal{R}^{d \times 1}, \quad \mathbf{y} \in \mathcal{R}^{m \times 1}$$

$$\mathbf{X} = \begin{bmatrix} x_{1,1} & x_{1,2} & \cdots & x_{1,d} \\ \vdots & \vdots & \ddots & \vdots \\ x_{m,1} & x_{m,2} & \cdots & x_{m,d} \end{bmatrix}, \quad \mathbf{w} = \begin{bmatrix} w_1 \\ \vdots \\ w_d \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_m \end{bmatrix}$$

Linear Regression

Problem Statement: To predict the unknown y for a given \mathbf{x} (**testing**)

- We have a collection of labeled examples (**training**) $\{(\mathbf{x}_i, y_i)\}_{i=1}^m$
 - m is the size of the collection
 - \mathbf{x}_i is the d -dimensional feature vector of example $i = 1, \dots, m$ (input)
 - y_i is a real-valued target (1-D)
 - Note:
 - when y_i is **continuous** valued, it is a **regression problem**
 - when y_i is **discrete** valued, it is a **classification problem**
- We want to **build a model** $f_{\mathbf{w},b}(\mathbf{x})$ as a linear combination of features of example \mathbf{x} : $f_{\mathbf{w},b}(\mathbf{x}) = \mathbf{x}^T \mathbf{w} + b$
 where \mathbf{w} is a d -dimensional vector of parameters and b is a real number.
- The notation $f_{\mathbf{w},b}$ means that the model f is parametrized by two values: \mathbf{w} and b

Ref: [Book4] Stephen Boyd and Lieven Vandenberghe, "Introduction to Applied Linear Algebra", Cambridge University Press, 2018 (chp.14)

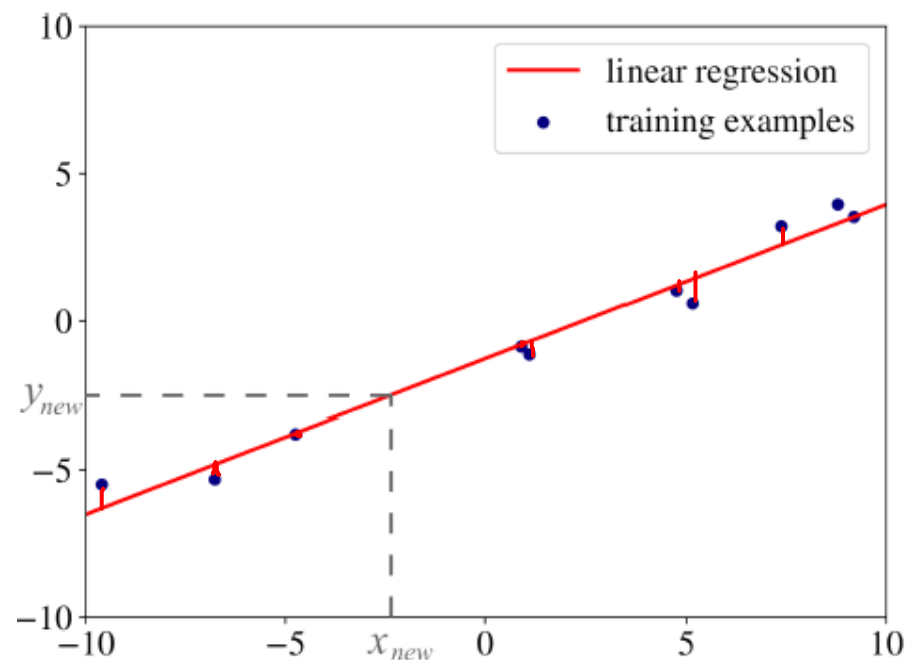
Linear Regression

Learning objective function

- To find the optimal values for \mathbf{w}^* and b^* which **minimizes** the following expression:

$$\frac{1}{m} \sum_{i=1}^m (f_{\mathbf{w},b}(\mathbf{x}_i) - y_i)^2$$

- In mathematics, the expression we minimize or maximize is called an **objective function**, or, simply, an **objective**



$(f_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2$ is called the **loss function**: a measure of the difference between $f_{\mathbf{w}}(\mathbf{x}_i)$ and y_i or a penalty for misclassification of example i .

Ref: [Book1] Andriy Burkov, "The Hundred-Page Machine Learning Book", 2019 (chp3.1.2)

Linear Regression

Learning objective function (using simplified notation hereon)

- To find the optimal values for \mathbf{w}^* which **minimizes** the following expression:

$$\sum_{i=1}^m (f_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2$$

with $f_{\mathbf{w}}(\mathbf{x}_i) = \mathbf{x}_i^T \mathbf{w}$,

where we define $\mathbf{w} = [b, w_1, \dots, w_d]^T = [w_0, w_1, \dots, w_d]^T$,

and $\mathbf{x}_i = [1, x_{i,1}, \dots, x_{i,d}]^T = [x_{i,0}, x_{i,1}, \dots, x_{i,d}]^T, i = 1, \dots, m$

- This particular choice of the loss function is called **squared error loss**

Note: The normalization factor $\frac{1}{m}$ can be omitted as it does not affect the optimization.

Linear Regression

$$\sum_{i=1}^m (f_{\mathbf{w},b}(\mathbf{x}_i) - y_i)^2$$

- All model-based learning algorithms have a **loss function**
- What we do to find the best model is **to minimize the objective** known as the **cost function**
- **Cost function** is a sum of **loss functions** over training set plus possibly some model complexity penalty (regularization)
- In linear regression, the cost function is given by the *average loss*, also called the **empirical risk** because we do not have all the data (e.g. testing data)
 - The average of all penalties is obtained by applying the model to the training data

Linear Regression

Learning (Training)

- Consider the set of feature vector \mathbf{x}_i and target output y_i indexed by $i = 1, \dots, m$, a linear model $f_{\mathbf{w}}(\mathbf{x}) = \mathbf{x}^T \mathbf{w}$ can be stacked as

$$\begin{aligned}
 & \text{Learning Model} \rightarrow f_{\mathbf{w}}(\mathbf{X}) = \mathbf{X}\mathbf{w} \quad \longleftrightarrow \quad \mathbf{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_m \end{bmatrix} \leftarrow \text{Learning target vector} \\
 & \quad \quad \quad = \begin{bmatrix} \mathbf{x}_1^T \mathbf{w} \\ \vdots \\ \mathbf{x}_m^T \mathbf{w} \end{bmatrix} \\
 & \text{where } \mathbf{x}_i^T \mathbf{w} = [1, x_{i,1}, \dots, x_{i,d}] \begin{bmatrix} b \\ w_1 \\ \vdots \\ w_d \end{bmatrix}
 \end{aligned}$$

Note: The **bias/offset term** is responsible for **translating** the line/plane/hyperplane away from the origin.

Linear Regression

Least Squares Regression

In vector-matrix notation, the minimization of the objective function can be written compactly using $\mathbf{e} = \mathbf{X}\mathbf{w} - \mathbf{y}$:

$$\begin{aligned} J(\mathbf{w}) &= \mathbf{e}^T \mathbf{e} \\ &= (\mathbf{X}\mathbf{w} - \mathbf{y})^T (\mathbf{X}\mathbf{w} - \mathbf{y}) \\ &= (\mathbf{w}^T \mathbf{X}^T - \mathbf{y}^T) (\mathbf{X}\mathbf{w} - \mathbf{y}) \\ &= \mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{w}^T \mathbf{X}^T \mathbf{y} - \mathbf{y}^T \mathbf{X} \mathbf{w} + \mathbf{y}^T \mathbf{y} \\ &= \mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - 2\mathbf{y}^T \mathbf{X} \mathbf{w} + \mathbf{y}^T \mathbf{y}. \end{aligned}$$

Note: when $f_{\mathbf{w}}(\mathbf{X}) = \mathbf{X}\mathbf{w}$, then

$$\sum_{i=1}^m (f_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2 = (\mathbf{X}\mathbf{w} - \mathbf{y})^T (\mathbf{X}\mathbf{w} - \mathbf{y}).$$

Linear Regression

Differentiating $J(\mathbf{w})$ with respect to \mathbf{w} and setting the result to $\mathbf{0}$:

$$\frac{\partial}{\partial \mathbf{w}} J(\mathbf{w}) = \mathbf{0}$$

$$\frac{\partial}{\partial \mathbf{w}} (\mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - 2 \mathbf{y}^T \mathbf{X} \mathbf{w} + \mathbf{y}^T \mathbf{y}) = \mathbf{0}$$

$$\Rightarrow 2 \mathbf{X}^T \mathbf{X} \mathbf{w} - 2 \mathbf{X}^T \mathbf{y} = \mathbf{0}$$

$$\Rightarrow \mathbf{X}^T \mathbf{X} \mathbf{w} = \mathbf{X}^T \mathbf{y}$$

\Rightarrow Any minimizer $\hat{\mathbf{w}}$ of $J(\mathbf{w})$ must satisfy $\mathbf{X}^T (\mathbf{X} \mathbf{w} - \mathbf{y}) = \mathbf{0}$.

If $\mathbf{X}^T \mathbf{X}$ is invertible, then

Learning/training:

$$\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Prediction/testing:

$$\hat{f}_{\mathbf{w}}(\mathbf{X}_{new}) = \mathbf{X}_{new} \hat{\mathbf{w}}$$

Linear Regression

Example 1 Training set $\{(x_i, y_i)\}_{i=1}^m$

$$\mathbf{X} \quad \mathbf{w} \quad \mathbf{y}$$

$$\begin{bmatrix} 1 & -9 \\ 1 & -7 \\ 1 & -5 \\ 1 & 1 \\ 1 & 5 \\ 1 & 9 \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} = \begin{bmatrix} -6 \\ -6 \\ -4 \\ -1 \\ 1 \\ 4 \end{bmatrix}$$

$$\begin{aligned} \{x = -9\} &\rightarrow \{y = -6\} \\ \{x = -7\} &\rightarrow \{y = -6\} \\ \{x = -5\} &\rightarrow \{y = -4\} \\ \{x = 1\} &\rightarrow \{y = -1\} \\ \{x = 5\} &\rightarrow \{y = 1\} \\ \{x = 9\} &\rightarrow \{y = 4\} \end{aligned}$$

This set of linear equations has **no exact solution**

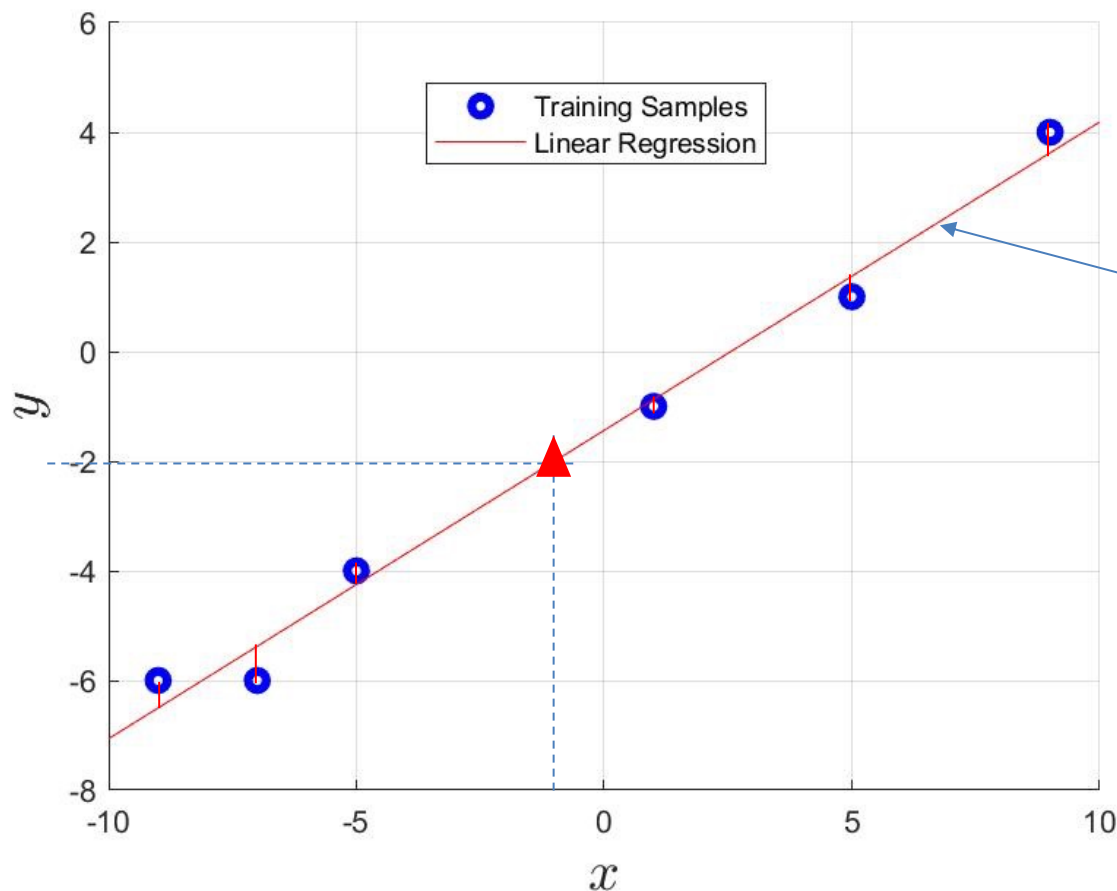
However, $\mathbf{X}^T \mathbf{X}$ is invertible

Least square approximation

$$\hat{\mathbf{w}} = \mathbf{X}^\dagger \mathbf{y} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

$$= \begin{bmatrix} 6 & -6 \\ -6 & 262 \end{bmatrix}^{-1} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 & 1 \\ -9 & -7 & -5 & 1 & 5 & 9 \end{bmatrix} \begin{bmatrix} -6 \\ -6 \\ -4 \\ -1 \\ 1 \\ 4 \end{bmatrix} = \begin{bmatrix} -1.4375 \\ 0.5625 \end{bmatrix}$$

Linear Regression



Linear Regression on one-dimensional samples

$$\begin{aligned}\hat{y} &= \mathbf{X}\hat{\mathbf{w}} \\ &= \mathbf{X} \begin{bmatrix} -1.4375 \\ 0.5625 \end{bmatrix} \\ y &= -1.4375 + 0.5625x\end{aligned}$$

Prediction:
Test set

$$\{x = -1\} \rightarrow \{y = ?\}$$

$$\begin{aligned}\hat{y} &= [1 \quad -1] \begin{bmatrix} -1.4375 \\ 0.5625 \end{bmatrix} \\ &= -2\end{aligned}$$

Python demo 1

Linear Regression

Example 2 $\{(\mathbf{x}_i, y_i)\}_{i=1}^m$

Training set

$$\{x_1 = 1, x_2 = 1, x_3 = 1\} \rightarrow \{y = 1\}$$

$$\{x_1 = 1, x_2 = -1, x_3 = 1\} \rightarrow \{y = 0\}$$

$$\{x_1 = 1, x_2 = 1, x_3 = 3\} \rightarrow \{y = 2\}$$

$$\{x_1 = 1, x_2 = 1, x_3 = 0\} \rightarrow \{y = -1\}$$

$$\mathbf{X} \quad \mathbf{w} \quad \mathbf{y}$$

$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & -1 & 1 \\ 1 & 1 & 3 \\ 1 & 1 & 0 \end{bmatrix} \begin{bmatrix} w_1 \\ w_2 \\ w_3 \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ 2 \\ -1 \end{bmatrix}$$

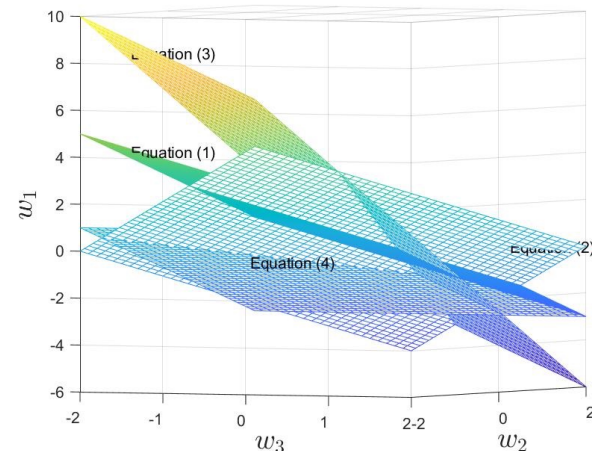
This set of linear equations has **no exact solution**

However, $\mathbf{X}^T \mathbf{X}$ is invertible

$$\hat{\mathbf{w}} = \mathbf{X}^\dagger \mathbf{y} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

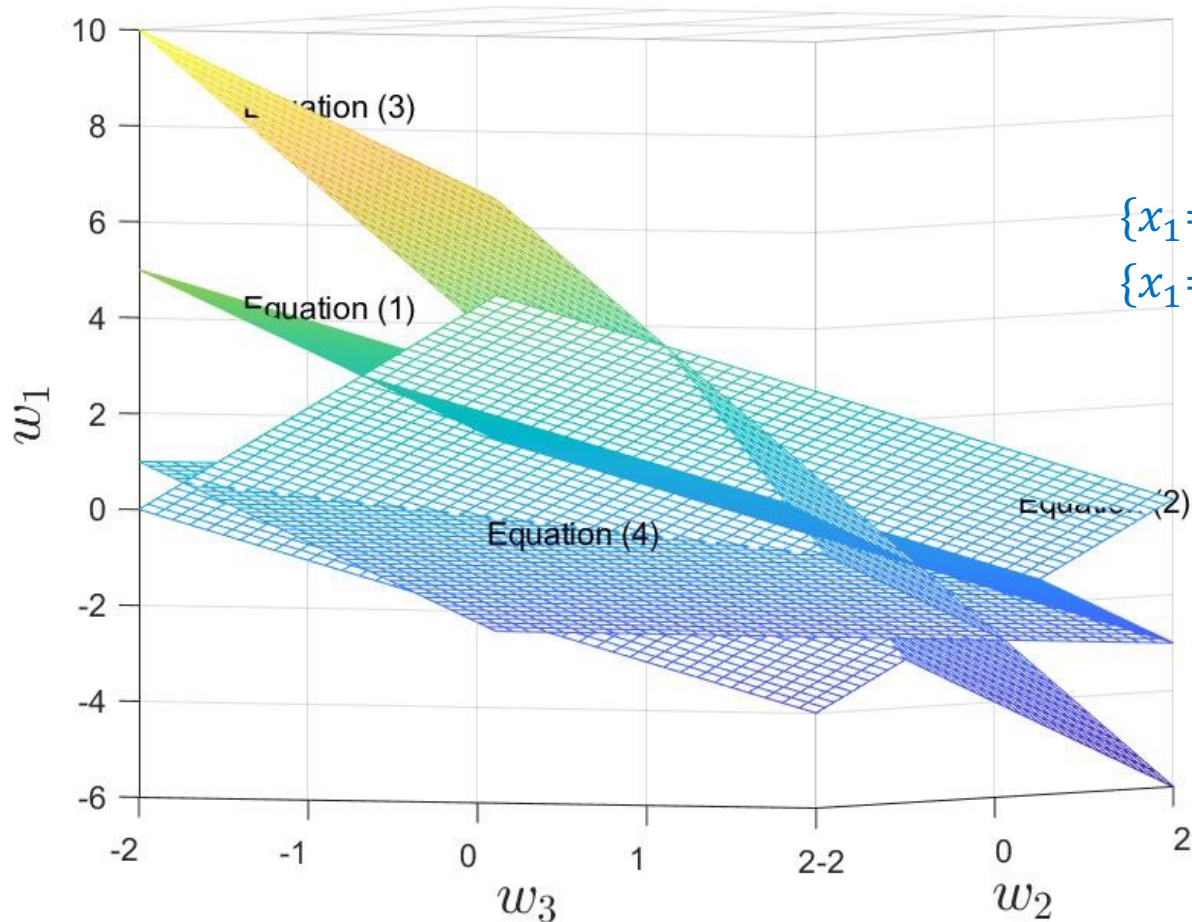
Least square approximation

$$= \begin{bmatrix} 4 & 2 & 5 \\ 2 & 4 & 3 \\ 5 & 3 & 11 \end{bmatrix}^{-1} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & -1 & 1 & 1 \\ 1 & 1 & 3 & 0 \end{bmatrix} \begin{bmatrix} 1 \\ 0 \\ 2 \\ -1 \end{bmatrix} = \begin{bmatrix} -0.7500 \\ 0.1786 \\ 0.9286 \end{bmatrix}$$



Linear Regression

The four linear equations



Prediction:

Test set

$$\{x_1 = 1, x_2 = 6, x_3 = 8\} \rightarrow \{y = ?\}$$

$$\{x_1 = 1, x_2 = 0, x_3 = -1\} \rightarrow \{y = ?\}$$

$$\hat{y} = \hat{f}_{\mathbf{w}}(\mathbf{X}_{new}) = \mathbf{X}_{new} \hat{\mathbf{w}}$$

$$\hat{y} = \begin{bmatrix} 1 & 6 & 8 \\ 1 & 0 & -1 \end{bmatrix} \begin{bmatrix} -0.7500 \\ 0.1786 \\ 0.9286 \end{bmatrix} = \begin{bmatrix} 7.7500 \\ -1.6786 \end{bmatrix}$$

Linear Regression

Learning of **Vectored Function** (Multiple Outputs)

For one sample: a linear model $\mathbf{f}_{\mathbf{w}}(\mathbf{x}) = \mathbf{x}^T \mathbf{W}$

Vector function

For m samples: $\mathbf{F}_{\mathbf{w}}(\mathbf{X}) = \mathbf{X}\mathbf{W} = \mathbf{Y}$

$$\begin{array}{l} \text{Sample 1} \cdots \rightarrow \\ \vdots \\ \text{Sample } m \cdots \rightarrow \end{array} \begin{bmatrix} \mathbf{x}_1^T \\ \vdots \\ \mathbf{x}_m^T \end{bmatrix} \mathbf{W} = \begin{bmatrix} 1 & x_{1,1} & \cdots & x_{1,d} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_{m,1} & \cdots & x_{m,d} \end{bmatrix} \underbrace{\begin{bmatrix} w_{0,1} & \cdots & w_{0,h} \\ w_{1,1} & \cdots & w_{1,h} \\ \vdots & \ddots & \vdots \\ w_{d,1} & \cdots & w_{d,h} \end{bmatrix}}_h$$

$$\begin{array}{l} \text{Sample 1's output} \cdots \rightarrow \\ \vdots \\ \text{Sample } m\text{'s output} \cdots \rightarrow \end{array} \begin{bmatrix} y_{1,1} & \cdots & y_{1,h} \\ \vdots & \vdots & \vdots \\ y_{m,1} & \cdots & y_{m,h} \end{bmatrix} \underbrace{\quad}_h \quad m$$

$$\mathbf{X} \in \mathcal{R}^{m \times (d+1)}, \mathbf{W} \in \mathcal{R}^{(d+1) \times h}, \mathbf{Y} \in \mathcal{R}^{m \times h}$$

Linear Regression

Objective: $\sum_{i=1}^m (\mathbf{f}_{\mathbf{w}}(\mathbf{x}_i) - \mathbf{y}_i)^2 = \mathbf{E}^T \mathbf{E}$

Least Squares Regression of Multiple Outputs

In matrix notation, the sum of squared errors cost function can be written compactly using $\mathbf{E} = \mathbf{XW} - \mathbf{Y}$:

$$\begin{aligned} J(\mathbf{W}) &= \text{trace}(\mathbf{E}^T \mathbf{E}) \\ &= \text{trace}[(\mathbf{XW} - \mathbf{Y})^T (\mathbf{XW} - \mathbf{Y})] \end{aligned}$$

If $\mathbf{X}^T \mathbf{X}$ is invertible, then

Learning/training: $\hat{\mathbf{W}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$

Prediction/testing: $\hat{\mathbf{F}}_{\mathbf{w}}(\mathbf{X}_{new}) = \mathbf{X}_{new} \hat{\mathbf{W}}$

Ref: Hastie, Tibshirani, Friedman, "The Elements of Statistical Learning", (2nd ed., 12th printing) 2017 (chp.3.2.4)

Least Squares Regression of Multiple Outputs

$$J(\mathbf{W}) = \text{trace}(\mathbf{E}^T \mathbf{E})$$

$$= \text{trace}\left(\begin{bmatrix} \mathbf{e}_1^T \\ \vdots \\ \mathbf{e}_h^T \end{bmatrix} [\mathbf{e}_1 \quad \mathbf{e}_2 \quad \dots \quad \mathbf{e}_h]\right)$$

$$= \text{trace}\left(\begin{bmatrix} \mathbf{e}_1^T \mathbf{e}_1 & \mathbf{e}_1^T \mathbf{e}_2 & \dots & \mathbf{e}_1^T \mathbf{e}_h \\ \mathbf{e}_2^T \mathbf{e}_1 & \mathbf{e}_2^T \mathbf{e}_2 & \dots & \mathbf{e}_2^T \mathbf{e}_h \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{e}_h^T \mathbf{e}_1 & \mathbf{e}_h^T \mathbf{e}_2 & \dots & \mathbf{e}_h^T \mathbf{e}_h \end{bmatrix}\right) = \sum_{k=1}^h \mathbf{e}_k^T \mathbf{e}_k$$

Linear Regression of multiple outputs

Example 3

Training set $\{x_1 = 1, x_2 = 1, x_3 = 1\} \rightarrow \{y_1 = 1, y_2 = 0\}$

$\{\mathbf{x}_i, \mathbf{y}_i\}_{i=1}^m$ $\{x_1 = 1, x_2 = -1, x_3 = 1\} \rightarrow \{y_1 = 0, y_2 = 1\}$

$\{x_1 = 1, x_2 = 1, x_3 = 3\} \rightarrow \{y_1 = 2, y_2 = -1\}$

$\{x_1 = 1, x_2 = 1, x_3 = 0\} \rightarrow \{y_1 = -1, y_2 = 3\}$

$$\begin{array}{c} \mathbf{X} \\ \text{Bias} \rightarrow \end{array} \begin{bmatrix} 1 & 1 & 1 \\ 1 & -1 & 1 \\ 1 & 1 & 3 \\ 1 & 1 & 0 \end{bmatrix} \begin{bmatrix} w_{1,1} & w_{1,2} \\ w_{2,1} & w_{2,2} \\ w_{3,1} & w_{3,2} \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 2 & -1 \\ -1 & 3 \end{bmatrix}$$

This set of linear equations has NO exact solution

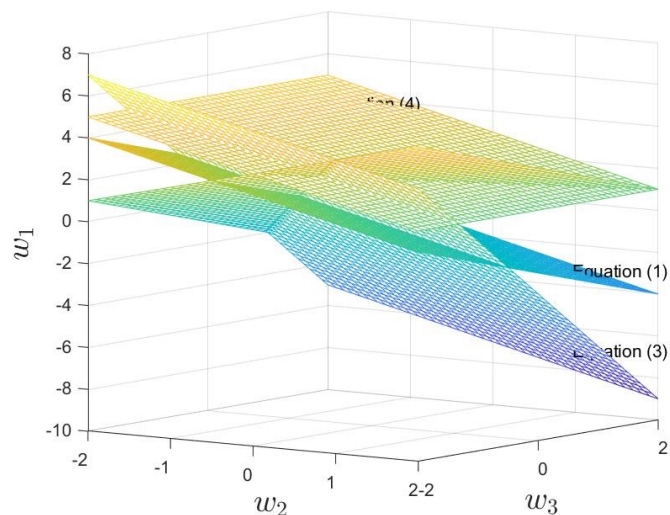
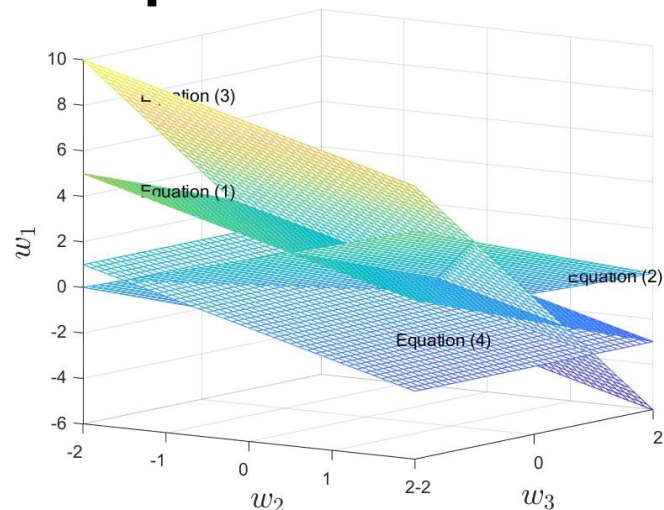
$$\hat{\mathbf{W}} = \mathbf{X}^{\dagger} \mathbf{Y} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y} \quad \mathbf{X}^T \mathbf{X} \text{ is invertible}$$

Least square approximation

$$= \begin{bmatrix} 4 & 2 & 5 \\ 2 & 4 & 3 \\ 5 & 3 & 11 \end{bmatrix}^{-1} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & -1 & 1 & 1 \\ 1 & 1 & 3 & 0 \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 2 & -1 \\ -1 & 3 \end{bmatrix} = \begin{bmatrix} -0.75 & 2.25 \\ 0.1786 & 0.0357 \\ 0.9286 & -1.2143 \end{bmatrix}$$

Linear Regression of multiple outputs

Example 3



Prediction:

Test set: two new samples

$$\{x_1 = 1, x_2 = 6, x_3 = 8\} \rightarrow \{y_1 = ?, y_2 = ?\}$$

$$\{x_1 = 1, x_2 = 0, x_3 = -1\} \rightarrow \{y_1 = ?, y_2 = ?\}$$

$$\hat{\mathbf{Y}} = \mathbf{X}_{new} \hat{\mathbf{W}}$$

$$\begin{aligned} \text{Bias} &\rightarrow \begin{bmatrix} 1 & 6 & 8 \\ 1 & 0 & -1 \end{bmatrix} \begin{bmatrix} -0.75 & 2.25 \\ 0.1786 & 0.0357 \\ 0.9286 & -1.2143 \end{bmatrix} \\ &= \begin{bmatrix} 7.75 & -7.25 \\ -1.6786 & 3.4643 \end{bmatrix} \end{aligned}$$

Python demo 2

Linear Regression of multiple outputs

Example 4

The values of feature x and their corresponding values of multiple outputs target y are shown in the table below.

Based on the least square regression, what are the values of \mathbf{w} ?
Based on the current mapping, when $x = 2$, what is the value of y ?

x	[3]	[4]	[10]	[6]	[7]
y	[0, 5]	[1.5, 4]	[-3, 8]	[-4, 10]	[1, 6]

$$\hat{\mathbf{W}} = \mathbf{X}^\dagger \mathbf{Y} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y} = \begin{bmatrix} 1.9 & 3.6 \\ -0.4667 & 0.5 \end{bmatrix}$$

$$\widehat{\mathbf{Y}}_{new} = \mathbf{X}_{new} \hat{\mathbf{W}} = \begin{bmatrix} 1 & 2 \end{bmatrix} \hat{\mathbf{W}} = \begin{bmatrix} 0.9667 & 4.6 \end{bmatrix}$$

Python demo 3

Prediction

Summary

Midterm (L1 to L5)
 Trial quiz

- Notations, Vectors, Matrices
- Operations on Vectors and Matrices
 - Dot-product, matrix inverse
- Systems of Linear Equations $f_{\mathbf{w}}(\mathbf{X}) = \mathbf{X}\mathbf{w} = \mathbf{y}$
 - Matrix-vector notation, linear dependency, invertible
 - Even-, over-, under-determined linear systems
- Functions, Derivative and Gradient
 - Inner product, linear/affine functions
 - Maximum and minimum, partial derivatives, gradient
- Least Squares, Linear Regression
 - Objective function, loss function
 - Least square solution, training/learning and testing/prediction
 - Linear regression with multiple outputs

Learning/training $\hat{\mathbf{w}} = (\mathbf{X}_{train}^T \mathbf{X}_{train})^{-1} \mathbf{X}_{train}^T \mathbf{y}_{train}$

Prediction/testing $\mathbf{y}_{test} = \mathbf{X}_{test} \hat{\mathbf{w}}$

- Classification
 - Ridge Regression
 - Polynomial Regression
- Python packages: numpy, pandas, matplotlib.pyplot, numpy.linalg, and sklearn.metrics (for mean_squared_error), numpy.linalg.pinv