

# EE2211 Introduction to Machine Learning

## Lecture 8

Yueming Jin

[ymjin@nus.edu.sg](mailto:ymjin@nus.edu.sg)

Electrical and Computer Engineering Department  
National University of Singapore

*Acknowledgement: EE2211 development team  
Thomas, Helen, Xinchao, Kar-Ann, Chen Khong, Robby and Haizhou*

# Course Contents

- Introduction and Preliminaries (Xinchao)
  - Introduction
  - Data Engineering
  - Introduction to Probability and Statistics
- Fundamental Machine Learning Algorithms I (Yueming)
  - Systems of linear equations
  - Least squares, Linear regression
  - Ridge regression, Polynomial regression
- Fundamental Machine Learning Algorithms II (Yueming)
  - Over-fitting, bias/variance trade-off
  - Optimization, Gradient descent
  - Decision Trees, Random Forest
- Performance and More Algorithms (Xinchao)
  - Performance Issues
  - K-means Clustering
  - Neural Networks

# Fundamental ML Algorithms: Optimization, Gradient Descent

## Module III Contents

- Overfitting, underfitting and model complexity
- Regularization
- Bias-variance trade-off
- Loss function
- Optimization
- Gradient descent
- Decision trees
- Random forest

# Review

- Supervised learning: given feature(s)  $x$ , we want to predict target  $y$
- Most supervised learning algorithms can be formulated as the following optimization problem

$$\operatorname{argmin}_{\mathbf{w}} \mathbf{Data-Loss}(\mathbf{w}) + \lambda \mathbf{Regularization}(\mathbf{w})$$

- **Data-Loss(w)** quantifies fitting error to training set given parameters  $\mathbf{w}$ : smaller error  $\Rightarrow$  better fit to training data
- **Regularization(w)** penalizes more complex models
- For example, in the case of polynomial regression (previous lectures):

$$\operatorname{argmin}_{\mathbf{w}} \underbrace{(\mathbf{P}\mathbf{w} - \mathbf{y})^T (\mathbf{P}\mathbf{w} - \mathbf{y})}_{\mathbf{Data-Loss}(\mathbf{w})} + \lambda \underbrace{\mathbf{w}^T \mathbf{w}}_{\mathbf{Reg}(\mathbf{w})}$$

# Review

- For polynomial regression (previous lectures)

$$\underset{\mathbf{w}}{\operatorname{argmin}} C(\mathbf{w}) = \underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{P}\mathbf{w} - \mathbf{y})^T (\mathbf{P}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w}$$

$$= \underset{\mathbf{w}}{\operatorname{argmin}} \sum_{i=1}^m (\mathbf{p}_i^T \mathbf{w} - y_i)^2 + \lambda \mathbf{w}^T \mathbf{w}$$

$\mathbf{p}_i^T \mathbf{w}$  is prediction of  $i$ -th  
training sample

$y_i$  is target of  $i$ -th  
training sample

# Review

- For polynomial regression (previous lectures)

$$\begin{aligned}\underset{\mathbf{w}}{\operatorname{argmin}} C(\mathbf{w}) &= \underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{P}\mathbf{w} - \mathbf{y})^T (\mathbf{P}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w} \\ &= \underset{\mathbf{w}}{\operatorname{argmin}} \sum_{i=1}^m (\mathbf{p}_i^T \mathbf{w} - y_i)^2 + \lambda \mathbf{w}^T \mathbf{w}\end{aligned}$$

- Linear regression with 2 features,  $\mathbf{p}_i = \begin{bmatrix} 1 \\ x_1 \\ x_2 \end{bmatrix}_i$ 
  - ← Bias/Offset
  - ← Feature 1 of i-th sample
  - ← Feature 2 of i-th sample
- Quadratic regression with 1 feature,  $\mathbf{p}_i = \begin{bmatrix} 1 \\ x \\ x^2 \end{bmatrix}_i$ 
  - ← Bias/Offset
  - ←  $x$  is feature of i-th sample
  - ←  $x^2$  is feature of i-th sample

# Loss Function & Learning Model

- For polynomial regression (previous lectures)

$$\begin{aligned}\operatorname{argmin}_{\mathbf{w}} C(\mathbf{w}) &= \operatorname{argmin}_{\mathbf{w}} (\mathbf{P}\mathbf{w} - \mathbf{y})^T (\mathbf{P}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w} \\ &= \operatorname{argmin}_{\mathbf{w}} \sum_{i=1}^m (\mathbf{p}_i^T \mathbf{w} - y_i)^2 + \lambda \mathbf{w}^T \mathbf{w}\end{aligned}$$

- Let  $f(\mathbf{x}_i, \mathbf{w})$  be the prediction of target  $y_i$  from features  $\mathbf{x}_i$  for  $i$ -th training sample. For example, suppose  $f(\mathbf{x}_i, \mathbf{w}) = \mathbf{p}_i^T \mathbf{w}$ , then above becomes

$$\operatorname{argmin}_{\mathbf{w}} C(\mathbf{w}) = \operatorname{argmin}_{\mathbf{w}} \sum_{i=1}^m (f(\mathbf{x}_i, \mathbf{w}) - y_i)^2 + \lambda \mathbf{w}^T \mathbf{w}$$

- Let  $L(f(\mathbf{x}_i, \mathbf{w}), y_i)$  be the penalty for predicting  $f(\mathbf{x}_i, \mathbf{w})$  when true value is  $y_i$ . For example, suppose  $L(f(\mathbf{x}_i, \mathbf{w}), y_i) = (f(\mathbf{x}_i, \mathbf{w}) - y_i)^2$ , then above becomes

$$\operatorname{argmin}_{\mathbf{w}} C(\mathbf{w}) = \operatorname{argmin}_{\mathbf{w}} \sum_{i=1}^m L(f(\mathbf{x}_i, \mathbf{w}), y_i) + \lambda \mathbf{w}^T \mathbf{w}$$

# Loss Function & Learning Model

- From previous slide

$$\operatorname{argmin}_{\mathbf{w}} C(\mathbf{w}) = \operatorname{argmin}_{\mathbf{w}} \sum_{i=1}^m L(f(\mathbf{x}_i, \mathbf{w}), y_i) + \lambda \mathbf{w}^T \mathbf{w}$$

- To make it even more general, we can write

$$\operatorname{argmin}_{\mathbf{w}} C(\mathbf{w}) = \operatorname{argmin}_{\mathbf{w}} \sum_{i=1}^m L(f(\mathbf{x}_i, \mathbf{w}), y_i) + \lambda R(\mathbf{w})$$

Cost Function

Loss Function

Learning Model

Regularization



# Building Blocks of ML algorithms

- From previous slide

$$\operatorname{argmin}_{\mathbf{w}} C(\mathbf{w}) = \operatorname{argmin}_{\mathbf{w}} \sum_{i=1}^m L(f(\mathbf{x}_i, \mathbf{w}), y_i) + \lambda \mathbf{w}^T \mathbf{w}$$

- To make it even more general, we can write

$$\operatorname{argmin}_{\mathbf{w}} C(\mathbf{w}) = \operatorname{argmin}_{\mathbf{w}} \sum_{i=1}^m L(f(\mathbf{x}_i, \mathbf{w}), y_i) + \lambda R(\mathbf{w})$$

- **Learning model**  $f$  reflects our belief about the relationship between the features  $\mathbf{x}_i$  & target  $y_i$
- **Loss function**  $L$  is the penalty for predicting  $f(\mathbf{x}_i, \mathbf{w})$  when the true value is  $y_i$
- **Regularization**  $R$  encourages less complex models
- **Cost function**  $C$  is the final optimization criterion we want to minimize
- **Optimization routine** to find solution to cost function

# Motivation for Gradient Descent

- Different learning function  $f$ , loss function  $L$  & regularization  $R$  give rise to different learning algorithms
- In polynomial regression (previous lectures), optimal  $\mathbf{w}$  can be written with the following “closed-form” formula (primal solution):

$$\hat{\mathbf{w}} = (\mathbf{P}_{train}^T \mathbf{P}_{train} + \lambda \mathbf{I})^{-1} \mathbf{P}_{train}^T \mathbf{y}_{train}$$

- For other learning function  $f$ , loss function  $L$  & regularization  $R$ , optimizing  $C(\mathbf{w})$  might not be so easy
- Usually have to estimate  $\mathbf{w}$  iteratively with some algorithm
- Optimization workhorse for modern machine learning is gradient descent

# Gradient Descent Algorithm

- Suppose we want to minimize  $C(\mathbf{w})$  with respect to  $\mathbf{w} = [w_1, \dots, w_d]^T$

- Gradient  $\nabla_{\mathbf{w}} C(\mathbf{w}) = \begin{pmatrix} \frac{\partial C}{\partial w_1} \\ \frac{\partial C}{\partial w_2} \\ \vdots \\ \frac{\partial C}{\partial w_d} \end{pmatrix}$

- $\nabla_{\mathbf{w}} C(\mathbf{w})$  is vector & function of  $\mathbf{w}$

- $\nabla_{\mathbf{w}} C(\mathbf{w})$  is direction at  $\mathbf{w}$  where  $C$  is increasing most rapidly, so  $-\nabla_{\mathbf{w}} C(\mathbf{w})$  is direction at  $\mathbf{w}$  where  $C$  is decreasing most rapidly

- Gradient Descent:

Initialize  $\mathbf{w}_0$  and learning rate  $\eta$ ;

**while** *true* **do**


    Compute  $\mathbf{w}_{k+1} \leftarrow \mathbf{w}_k - \eta \nabla_{\mathbf{w}} C(\mathbf{w}_k)$

**if** *converge* **then**

**return**  $\mathbf{w}_{k+1}$

**end**

**end**



According to multi-variable calculus, if eta is not too big, then  $C(\mathbf{w}_{k+1}) < C(\mathbf{w}_k) \Rightarrow$  we get better  $\mathbf{w}$  after each iteration

# Gradient Descent Algorithm

- Gradient Descent:

Initialize  $\mathbf{w}_0$  and learning rate  $\eta$ ;

**while** *true* **do**

    Compute  $\mathbf{w}_{k+1} \leftarrow \mathbf{w}_k - \eta \nabla_{\mathbf{w}} C(\mathbf{w}_k)$

**if** *converge* **then**

**return**  $\mathbf{w}_{k+1}$

**end**

**end**


- Possible convergence criteria

- Set maximum iteration  $k$
- Check percentage or absolute change in  $C$  below a threshold
- Check percentage or absolute change in  $\mathbf{w}$  below a threshold

- Gradient descent can only find local minimum


- Because gradient = 0 at local minimum, so  $\mathbf{w}$  won't change after that

- Many variations of gradient descent, e.g., change how gradient is computed or learning rate  $\eta$  decreases with increasing  $k$




# Find $x$ to minimize $g(x) = x^2$

- Obviously minimum corresponds to  $x = 0$ , but let's do gradient descent
  - Gradient  $\nabla_x g(x) = 2x$
  - Initialize  $x_0 = -1$ , learning rate  $\eta = 0.4$
  - At each iteration,  $x_{k+1} = x_k - \eta \nabla_x g(x_k)$




# Find $x$ to minimize $g(x) = x^2$

- Obviously minimum corresponds to  $x = 0$ , but let's do gradient descent
  - Gradient  $\nabla_x g(x) = 2x$
  - Initialize  $x_0 = -1$ , learning rate  $\eta = 0.4$
  - At each iteration,  $x_{k+1} = x_k - \eta \nabla_x g(x_k)$




# Find $x$ to minimize $g(x) = x^2$

- Obviously minimum corresponds to  $x = 0$ , but let's do gradient descent
  - Gradient  $\nabla_x g(x) = 2x$
  - Initialize  $x_0 = -1$ , learning rate  $\eta = 0.4$
  - At each iteration,  $x_{k+1} = x_k - \eta \nabla_x g(x_k)$



# Find $x$ to minimize $g(x) = x^2$


- Obviously minimum corresponds to  $x = 0$ , but let's do gradient descent
  - Gradient  $\nabla_x g(x) = 2x$
  - Initialize  $x_0 = -1$ , learning rate  $\eta = 0.4$
  - At each iteration,  $x_{k+1} = x_k - \eta \nabla_x g(x_k)$






# Find $x$ to minimize $g(x) = x^2$

- Obviously minimum corresponds to  $x = 0$ , but let's do gradient descent
  - Gradient  $\nabla_x g(x) = 2x$
  - Initialize  $x_0 = -1$ , learning rate  $\eta = 0.4$
  - At each iteration,  $x_{k+1} = x_k - \eta \nabla_x g(x_k)$




# Find $x$ to minimize $g(x) = x^2$

- Obviously minimum corresponds to  $x = 0$ , but let's do gradient descent
  - Gradient  $\nabla_x g(x) = 2x$
  - Initialize  $x_0 = -1$ , learning rate  $\eta = 0.4$
  - At each iteration,  $x_{k+1} = x_k - \eta \nabla_x g(x_k)$




# What if learning rate $\eta$ is too small?

- Obviously minimum corresponds to  $x = 0$ , but let's do gradient descent
  - Gradient  $\nabla_x g(x) = 2x$
  - Initialize  $x_0 = -1$ , learning rate  $\eta = 0.01$
  - At each iteration,  $x_{k+1} = x_k - \eta \nabla_x g(x_k)$




# What if learning rate $\eta$ is too small?

- Obviously minimum corresponds to  $x = 0$ , but let's do gradient descent
  - Gradient  $\nabla_x g(x) = 2x$
  - Initialize  $x_0 = -1$ , learning rate  $\eta = 0.01$
  - At each iteration,  $x_{k+1} = x_k - \eta \nabla_x g(x_k)$




# What if learning rate $\eta$ is too small?

- Obviously minimum corresponds to  $x = 0$ , but let's do gradient descent
  - Gradient  $\nabla_x g(x) = 2x$
  - Initialize  $x_0 = -1$ , learning rate  $\eta = 0.01$
  - At each iteration,  $x_{k+1} = x_k - \eta \nabla_x g(x_k)$




# What if learning rate $\eta$ is too small?

- Obviously minimum corresponds to  $x = 0$ , but let's do gradient descent
  - Gradient  $\nabla_x g(x) = 2x$
  - Initialize  $x_0 = -1$ , learning rate  $\eta = 0.01$
  - At each iteration,  $x_{k+1} = x_k - \eta \nabla_x g(x_k)$




# What if learning rate $\eta$ is too small?

- Obviously minimum corresponds to  $x = 0$ , but let's do gradient descent
  - Gradient  $\nabla_x g(x) = 2x$
  - Initialize  $x_0 = -1$ , learning rate  $\eta = 0.01$
  - At each iteration,  $x_{k+1} = x_k - \eta \nabla_x g(x_k)$



# What if learning rate $\eta$ is too small?


- Obviously minimum corresponds to  $x = 0$ , but let's do gradient descent
  - Gradient  $\nabla_x g(x) = 2x$
  - Initialize  $x_0 = -1$ , learning rate  $\eta = 0.01$
  - At each iteration,  $x_{k+1} = x_k - \eta \nabla_x g(x_k)$






# What if learning rate $\eta$ is too big?

- Obviously minimum corresponds to  $x = 0$ , but let's do gradient descent
  - Gradient  $\nabla_x g(x) = 2x$
  - Initialize  $x_0 = -1$ , learning rate  $\eta = 0.9$
  - At each iteration,  $x_{k+1} = x_k - \eta \nabla_x g(x_k)$




# What if learning rate $\eta$ is too big?

- Obviously minimum corresponds to  $x = 0$ , but let's do gradient descent
  - Gradient  $\nabla_x g(x) = 2x$
  - Initialize  $x_0 = -1$ , learning rate  $\eta = 0.9$
  - At each iteration,  $x_{k+1} = x_k - \eta \nabla_x g(x_k)$




# What if learning rate $\eta$ is too big?

- Obviously minimum corresponds to  $x = 0$ , but let's do gradient descent
  - Gradient  $\nabla_x g(x) = 2x$
  - Initialize  $x_0 = -1$ , learning rate  $\eta = 0.9$
  - At each iteration,  $x_{k+1} = x_k - \eta \nabla_x g(x_k)$




# What if learning rate $\eta$ is too big?

- Obviously minimum corresponds to  $x = 0$ , but let's do gradient descent
  - Gradient  $\nabla_x g(x) = 2x$
  - Initialize  $x_0 = -1$ , learning rate  $\eta = 0.9$
  - At each iteration,  $x_{k+1} = x_k - \eta \nabla_x g(x_k)$



# What if learning rate $\eta$ is too big?


- Obviously minimum corresponds to  $x = 0$ , but let's do gradient descent
  - Gradient  $\nabla_x g(x) = 2x$
  - Initialize  $x_0 = -1$ , learning rate  $\eta = 0.9$
  - At each iteration,  $x_{k+1} = x_k - \eta \nabla_x g(x_k)$




# What if learning rate $\eta$ is too big?

- Obviously minimum corresponds to  $x = 0$ , but let's do gradient descent
  - Gradient  $\nabla_x g(x) = 2x$
  - Initialize  $x_0 = -1$ , learning rate  $\eta = 0.9$
  - At each iteration,  $x_{k+1} = x_k - \eta \nabla_x g(x_k)$

Python  
demo



# Gradient descent: quadratic function




Convergence to the foot of the valley

# Different Learning Models

- Different learning models  $f(\mathbf{x}_i, \mathbf{w})$  reflect our beliefs about the relationship between the features  $\mathbf{x}_i$  and target  $y_i$ 
  - For example,  $f(\mathbf{x}_i, \mathbf{w}) = \mathbf{p}_i^T \mathbf{w}$  assumes polynomial relationship between features and target
- Suppose we are performing classification (rather than regression), so  $y_i$  is class  $-1$  or class  $1$ 
  - $\mathbf{p}_i^T \mathbf{w}$  is number between  $-\infty$  to  $\infty$ .
  - Can use sigmoid function to map  $\mathbf{p}_i^T \mathbf{w}$  to between 0 and 1:

$$f(\mathbf{x}_i, \mathbf{w}) = \sigma(\mathbf{p}_i^T \mathbf{w})$$

$$\sigma(a) = \frac{1}{1 + e^{-a}}$$





# Different Learning Models

- Different learning models  $f(\mathbf{x}_i, \mathbf{w})$  reflect our beliefs about the relationship between the features  $\mathbf{x}_i$  and target  $y_i$ 
  - For example,  $f(\mathbf{x}_i, \mathbf{w}) = \mathbf{p}_i^T \mathbf{w}$  assumes polynomial relationship between features and target
- Suppose we are performing classification (rather than regression), so  $y_i$  is class  $-1$  or class  $1$ 
  - $\mathbf{p}_i^T \mathbf{w}$  is number between  $-\infty$  to  $\infty$ .
  - Can use sigmoid function to map  $\mathbf{p}_i^T \mathbf{w}$  to between  $0$  and  $1$ :



$$f(\mathbf{x}_i, \mathbf{w}) = \sigma(\mathbf{p}_i^T \mathbf{w})$$

$$\sigma(a) = \frac{1}{1 + e^{-a}}$$

- If  $f(\mathbf{x}_i, \mathbf{w})$  is closer to  $0$  (or  $1$ ), we predict class  $-1$  (or class  $1$ )
- More generally, in one layer neural network:  $f(\mathbf{x}_i, \mathbf{w}) = \sigma(\mathbf{p}_i^T \mathbf{w})$ , where activation function  $\sigma$  can be sigmoid or some other functions &  $\mathbf{p}$  is linear


# Different Learning Models

- $f(\mathbf{x}_i, \mathbf{w}) = \sigma(\mathbf{p}_i^T \mathbf{w})$ , where  $\sigma$  can be different functions:
- Rectified linear unit (ReLU):  $\sigma(a) = \max(0, a)$
- Exponential:  $\sigma(a) = \exp(-a)$



# Different Loss Functions

- Different loss functions  $L(f(\mathbf{x}_i, \mathbf{w}), y_i)$  encodes the penalty when we predict  $f(\mathbf{x}_i, \mathbf{w})$  but the true value is  $y_i$ 
  - $L(f(\mathbf{x}_i, \mathbf{w}), y_i) = (f(\mathbf{x}_i, \mathbf{w}) - y_i)^2$  is called the square error loss




# Different Loss Functions

- Different loss functions  $L(f(\mathbf{x}_i, \mathbf{w}), y_i)$  encodes the penalty when we predict  $f(\mathbf{x}_i, \mathbf{w})$  but the true value is  $y_i$ 
  - $L(f(\mathbf{x}_i, \mathbf{w}), y_i) = (f(\mathbf{x}_i, \mathbf{w}) - y_i)^2$  is called the square error loss
- Suppose we are performing classification (rather than regression), so  $y_i$  is class  $-1$  or class  $1$ , then square error loss makes less sense. Instead, we can use
  - Binary loss (or 0–1 loss):  $L(f(\mathbf{x}_i, \mathbf{w}), y_i) = \begin{cases} 0 & \text{if } f(\mathbf{x}_i, \mathbf{w}) = y_i \\ 1 & \text{if } f(\mathbf{x}_i, \mathbf{w}) \neq y_i \end{cases}$
  - In practice, hard to constrain  $f(\mathbf{x}_i, \mathbf{w})$  to be exactly  $-1$  or  $1$ , so we can declare “victory” if  $f(\mathbf{x}_i, \mathbf{w})$  &  $y$  have the same sign:

$$L(f(\mathbf{x}_i, \mathbf{w}), y_i) = \begin{cases} 0 & \text{if } f(\mathbf{x}_i, \mathbf{w})y_i > 0 \\ 1 & \text{if } f(\mathbf{x}_i, \mathbf{w})y_i < 0 \end{cases}$$

# Different Loss Functions

- Binary loss, where  $y_i$  is class  $-1$  or class  $1$  &  $f(\mathbf{x}_i, \mathbf{w})$  is a number between  $-\infty$  and  $\infty$ : 
$$L(f(\mathbf{x}_i, \mathbf{w}), y_i) = \begin{cases} 0 & \text{if } f(\mathbf{x}_i, \mathbf{w})y_i > 0 \\ 1 & \text{if } f(\mathbf{x}_i, \mathbf{w})y_i < 0 \end{cases}$$
- Binary loss not differentiable, so two other possibilities
  - Hinge loss:  $L(f(\mathbf{x}_i, \mathbf{w}), y_i) = \max(0, 1 - f(\mathbf{x}_i, \mathbf{w})y_i)$
  - Exponential loss:  $L(f(\mathbf{x}_i, \mathbf{w}), y_i) = \exp(-f(\mathbf{x}_i, \mathbf{w})y_i)$



# Summary

- Building blocks of machine learning algorithms
  - Learning model: reflects our belief about relationship between features & target we want to predict
  - Loss function: penalty for wrong prediction
  - Regularization: penalizes complex models
  - Optimization routine: find minimum of overall cost function
- Gradient descent algorithm
  - At each iteration, compute gradient & update model parameters in direction opposite to gradient
  - If learning rate  $\eta$  is too big  $\Rightarrow$  may not converge
  - If learning rate  $\eta$  is too small  $\Rightarrow$  converge very slowly
- Different learning models, e.g., linear, polynomial, sigmoid, ReLU, exponential, etc
- Different loss functions, e.g., square error, binary, logistic, etc